

Thinking...

Considering document structure

I'm looking at the v3 Table of Contents, which has 25 sections and no errors right now. I wonder if I should insert an Error catalogue as section 20 and then shift the validation section to 21. But, I'm questioning whether this would lead to any loss of clarity or if it would make maintaining the numbering easier. It's definitely worth exploring how to keep everything organized without losing meaning.

Considering literature audit

I'm thinking about preserving the historical claim since the prior draft marked certain elements, but the current version doesn't rely on that marker. I wonder if the user expects us to perform a literature audit based on their improvement suggestion. That might involve including a citation gate. Maybe I can browse to verify stable sources that reflect any recent changes or confirmations needed for accuracy. It feels like a good approach for solidifying our claims!

Clarifying document details

It seems I need to address the current date correction for the draft, which shows an incorrect future date of Aug 31. Also, v3a, dated Aug 24, indicates that all source verifications are still pending, unless there's supporting evidence. Notably, there's no runnable artifact, and I didn't execute a literature audit this time. While I could mention the lack of browsing due to this being a revision, I'm unsure if it's necessary to include that in the document. However, honesty remains important to me!

Clarifying error catalogues

I'm noticing a potential issue where the previous score suggested "restore consolidated error catalogue," but the original summary mentioned an expanded error catalogue, which isn't present. I need to add that detail to ensure clarity moving forward. It's important that my summaries accurately reflect the intended message, so I'll make sure to address this discrepancy. Let's see how we can rectify this to maintain consistency in our documentation!

common_agent_structure.v3a.md

Delivery note — read first.

Four constraints shape this revision:

1. **No runnable artifact was supplied.** There is no repository, harness implementation, model endpoint, collector, plugin runtime, memory backend, or accelerator in scope. Therefore v3a does **not** report **MEASURED_LOCAL** results. It delivers the complete architecture, quantitative release gates, executable harness contract, migration profile, consolidated error catalogue, and an honest separation of static, external, and unexecuted evidence.
2. **The supplied v3 draft was future-dated.** It carried a document date and research cutoff of **2026-08-31**, while this revision is issued on **2026-08-24**. v3a removes that future-dated status. No search, verification, measurement, or event after August 24, 2026 is represented as completed.

- 3. **No committed citation-audit artifact accompanied the supplied drafts.** References marked [V] by the unreleased v3 draft are therefore reclassified as [D]: *claimed verified by the prior draft, but not accepted as audit-verified by v3a*. CIT-GATE-001 blocks release until every [D], [C], and [K] reference is verified and the audit artifact is committed.
- 4. **v3a is standalone.** Requirements previously incorporated only by phrases such as “retained from v1/v2” are reproduced here. A consolidated error catalogue is restored. Fixed sample sizes are replaced by prospective power calculations, non-inferiority is separated from equivalence testing, and optional optimizer kill switches are distinguished from mandatory-control containment stops.

No v3 functional requirement has been intentionally removed. Where v3 contained an inconsistency, v3a retains the underlying control and supplies corrected wording.

Document ID: CASOPS-FS-COMMON-AGENT-STRUCTURE-V3A

Date: 2026-08-24

Status: Draft production implementation specification — release, deployment, self-improvement activation, and capability expansion remain human-gated and validation-blocked

Supersedes when approved: common_agent_structure.v2.md

Revises: unreleased common_agent_structure.v3.md draft labelled 2026-08-31

Host: common-agent-swarm-ops

Structure family: casops.common_agent.v3

Schema version: 3.0 — v3a is a specification correction and does not require a schema-family bump

Compatibility: v2 folders load through §22; v1 folders load through the v2 profile chained through §22

Research cutoff: 2026-08-24

Citation-audit status: BLOCKED

A v3a common agent remains **one self-contained folder and one agent_id**. It has nine first-class planes:

- 1. execution;
- 2. compatibility and protocol;
- 3. observability and decision provenance;
- 4. extensibility and plugins;
- 5. long-term memory;
- 6. autonomous improvement;
- 7. safety and adversarial robustness;
- 8. cache and context lifecycle;
- 9. corrigibility.

Table of contents

- 1. Purpose, changes, and defect register
- 2. Research basis, evidence policy, and citation audit
- 3. Core principles
- 4. Normative architecture
- 5. Folder contract
- 6. Composition and inheritance
- 7. Performance execution plane

8. Cache and context-lifecycle plane
 9. Compatibility and protocol plane
 10. Observability and decision provenance
 11. Extensible plugin architecture
 12. Long-term memory architecture
 13. Autonomous self-improvement
 14. Safety and adversarial-robustness plane
 15. Corrigibility plane
 16. Skills, identity, and persona isolation
 17. Compose and runtime algorithm
 18. Data models
 19. Operator and host APIs
 20. Consolidated error catalogue
 21. Validation specification, harness, and report
 22. Migration from v2
 23. Traceability
 24. Open risks
 25. Research references and citation audit
 26. Document control
-

1. Purpose, changes, and defect register

1.1 Purpose

v3a preserves the v1/v2 identity, inheritance, skill, disclosure, and fail-closed contracts while providing a complete production implementation specification for:

- task-completion latency, success, cost, and goodput;
- model, tool, protocol, and telemetry interoperability;
- end-to-end operational provenance;
- sandboxed executable extensions;
- typed, persistent, security-governed memory;
- gated autonomous improvement;
- cache and context lifecycle management;
- measured adversarial robustness;
- architecturally enforced corrigibility;
- powered and pre-registered validation.

The central validation thesis remains:

A specification must define not only what to measure, but how to measure it credibly.

The earlier v2 rule of “at least 30 observations” was inadequate for percentile and binary-rate gates. v3a requires prospective power calculations, paired or blocked task assignment, interval estimates, explicit estimands, and distinct procedures for superiority, non-inferiority, and equivalence.

1.2 Material changes

v2 to v3 architecture

Domain	v2	v3a
Performance	DAG, routing, opportunistic cache	Agent-aware cache scheduling, SLO admission control, goodput/CPE/CRR, adaptive compute and stopping
Context	Static segment budgets	Compaction, offload, re-grounding, sub-agent isolation, context-rot testing
Compatibility	Declared capability matrix	Asserted-versus-verified capabilities, capability-drift detection, tokenizer/template pinning, schema-profile negotiation
Observability	Traces, decision records, replay	Claim-level evidence graph, bounded reasoning monitor, mandatory tail retention, trace-cost budgets, automated RCA
Extensibility	Isolated process, digest, signature	I0–I3 isolation, object-capability handles, SBOM, provenance, scanning, ABI contracts
Memory	Seven stores, bitemporal records, hybrid retrieval	Paged hierarchy, trust tiers, poisoning resistance, consolidation, verified deletion and unlearning probes
Improvement	Propose-only candidates and human gate	Verifier-first objectives, reward-hacking detection, full regression ratchet, group-sequential canaries, immutable ledger
Safety	Distributed rules	First-class threat model, taint propagation, injection/hijack/exfiltration and termination gates
Corrigibility	Policy-based human gates	Host-owned invariant set unreachable by the agent
Validation	Fixed low sample floor and point comparisons	Prospective power, paired design, CIs, superiority/NI/equivalence separation, pre-registration
Citations	Inline references	Confidence markers and release-blocking audit

Unreleased v3 draft to v3a

Area	v3 draft issue	v3a correction
Date integrity	Dated 2026-08-31 before that date occurred	Dated 2026-08-24; no future completion claims
Citation state	Prior [V] claims had no attached audit artifact	Reclassified to [D] until audit evidence is committed
Binary sample sizes	$n \geq 400$ implied universal power for a 5pp gate	$n = \max(\text{floor}, \text{prospectively powered } n)$ using paired discordance assumptions
Statistical terminology	TOST described as the generic NI procedure	One-sided NI testing for NI; TOST only for two-sided equivalence

Area	v3 draft issue	v3a correction
Kill switches	"Every plane disableable" conflicted with mandatory safety/corrigibility	Optional optimizers have baseline kill switches; mandatory controls have fail-stop containment
Standalone use	Multiple requirements incorporated only by reference to v1/v2	Retained requirements reproduced
Error handling	Summary promised an expanded catalogue, but no consolidated section existed	Complete §20 added
Citation counts	Prior draft reported grouped counts not backed by a machine audit	Counts treated as provisional until <code>citation-audit.json</code> normalizes entries

1.3 Non-goals

This specification does not:

- authorize production activation;
- authorize unrestricted network access;
- authorize an agent to grant itself tools, plugins, permissions, or credentials;
- require or export private model chain-of-thought;
- permit the agent to modify its gates, telemetry, safety policy, permissions, or termination conditions;
- treat any preprint as a production baseline without local validation;
- claim bitwise model reproducibility without verified batch-invariant execution;
- claim that deletion from a memory store removes influence already incorporated into model weights;
- create a second public control plane;
- claim that external research gains are CASOPS gains;
- certify an implementation that was not supplied or executed.

1.4 Defect register

ID	Defect	Severity	v3a correction
DEF-001	v2 treated OpenTelemetry GenAI semantic conventions as a stable default contract.	High	GenAI attributes are treated as experimental until audit verifies otherwise. Pin <code>schema_url</code> , emit stable <code>casops.*</code> aliases, and treat any change as a compatibility event.
DEF-002	v2 labelled arXiv 2608.14624 as "CacheScout" and assigned an incorrect month/title.	Medium	The disputed label is removed. The identifier and title must be resolved by CIT-GATE-001 before use.
DEF-003	v2 attached a specific SWE-bench gain to an unverifiable Agent Lightning identifier.	High	The numeric claim remains withdrawn. Execution/training separation is retained on operational grounds and subject to citation audit.

ID	Defect	Severity	v3a correction
DEF-004	v2 used roughly 30 observations for p95 and binary success gates.	High	Replaced by powered, paired validation with minimum percentile floors and prospective binary-rate power calculations.
DEF-005	The supplied v3 draft was dated seven days in the future.	High	v3a uses August 24, 2026 and forbids representing future verification as complete.
DEF-006	v3 treated $n \geq 400$ as generally adequate for a 5pp binary gate and used TOST terminology for NI.	High	Binary n is power-derived from baseline and paired discordance; NI and equivalence procedures are separated in §21.4.
DEF-007	v3 said every added plane had a kill switch while declaring safety and corrigibility non-disableable.	Medium	Optional optimizers use kill switches; mandatory controls use containment stops that halt execution rather than bypass controls.
DEF-008	v3 depended on retained v1/v2 requirements without reproducing them and omitted a consolidated error catalogue.	Medium	Retained normative requirements are inlined and §20 is added.

2. Research basis, evidence policy, and citation audit

2.1 Search provenance

The supplied drafts report searches across:

- arXiv;
- NeurIPS, ICML, ICLR, PMLR, OpenReview, and ACL Anthology;
- W3C;
- CNCF;
- Linux Foundation;
- official OpenTelemetry, MCP, A2A, and CloudEvents material.

No machine-readable audit proving those searches was attached. v3a preserves the references and architecture but does not accept historical “verified this session” labels as release evidence.

2.2 Evidence maturity

Grade	Meaning	Production treatment
E1	Stable standard or peer-reviewed result with released evaluation	May default on only after local validation
E2	Peer-reviewed but workload-dependent, or an experimental standard	Feature-gated until local CASOPS gates pass

Grade	Meaning	Production treatment
E3	Recent preprint or early implementation	Experimental only, explicit flag, fallback, kill switch
E4	Open-ended self-modification or insufficiently bounded mechanism	Research-only, isolated, disabled in production

E-RULE-01. A feature supported only by E3 evidence must have:

- a validated fallback;
- an independently exercisable optimizer kill switch;
- a runtime budget;
- telemetry;
- no mandatory role on a high-risk execution path.

E-RULE-02. External reported deltas:

- are never additive;
- are never represented as CASOPS results;
- must identify the study or be labelled **TARGET**;
- cannot replace local validation.

E-RULE-03. Mandatory safety, audit, and corrigibility controls do not receive bypass kill switches. If unavailable, the system enters a containment stop.

2.3 Citation-confidence markers

Marker	Meaning	Release treatment
[A]	Accepted by a committed v3a citation-audit artifact	May support requirements subject to evidence grade
[D]	Prior unreleased v3 draft claimed verification; v3a has not accepted the audit	Blocked
[C]	Carried from v2 without a v3a audit	Blocked
[K]	Model or author knowledge without a v3a audit	Blocked

The prior draft reported a **14 / 31 / 13** split across verified, carried, and knowledge-derived references. Because the entries were grouped inconsistently and no audit artifact was supplied, those counts are historical metadata, not an accepted v3a inventory.

2.4 Evidence basis retained from the supplied draft

Serving and cache

The draft identifies work on:

- learned agent-execution-aware KV-cache management;
- policy-driven agentic serving;

- workload-aware multi-agent caching.

These motivate §8 but remain E3 and [D] until audited. Cache correctness, privacy, and fallback requirements are independently justified by operational engineering.

Memory

The draft identifies surveys of:

- agent memory mechanisms;
- memory evaluation limitations;
- long-term memory security;
- self-evolving and lifelong agents.

These motivate trust tiers, poisoning tests, evaluation-validity checks, and deletion verification. The controls remain defensible independently of any one citation.

Observability and compatibility

The draft reports:

- consolidation around OpenTelemetry-based GenAI observability;
- continued GenAI semantic-convention churn;
- rapid explicit MCP revisioning.

v3a therefore pins versions and emits a CASOPS-owned alias layer without assuming the current stability status of any unaudited external convention.

Self-improvement

The draft organizes improvement by what, when, how, and where evolution occurs. v3a retains that framing while enforcing the decisive boundary: prompts, context, memory, and workflow may produce candidates; permissions, gates, telemetry, safety, and production activation remain unreachable.

2.5 Citation audit gate

CIT-GATE-001 — release-blocking.

Before merge to `main`, every [D], [C], and [K] entry must be resolved to a live source with:

- matching identifier;
- matching title;
- authors;
- venue or repository;
- publication or revision date;
- evidence grade;
- located support for every attached numeric claim;
- audit actor;
- verification timestamp not later than the audit date;
- source digest where legally and technically available.

Output:

```
evals/reports/<run-id>/citation-audit.json
```

Minimum fields:

```
{
  "reference_id": "ref-001",
  "marker_before": "D",
  "resolved_identifier": "arXiv:...",
  "expected_title": "...",
  "observed_title": "...",
  "venue": "...",
  "year": 2026,
  "numeric_claims": [
    {
      "claim": "...",
      "location": "page/section/table",
      "status": "verified"
    }
  ],
  "verified_by": "reviewer-id",
  "verified_at": "2026-08-24T00:00:00Z",
  "status": "accepted"
}
```

Unresolvable references are deleted. Requirements depending solely on them must be independently justified or removed.

2.6 Date-integrity rule

CIT-GATE-002. A source, revision, benchmark result, or verification dated after the document’s audit date must not be represented as completed. Future-dated references may appear only as planned work and cannot support a requirement.

3. Core principles


ID	Principle	Normative meaning
P1	One identity	One folder remains one agent_id.
P2	Child owns mission	Parent content supports but never replaces the child mission.
P3	Safety tightens	Deny lists union, budgets take minima, and false safety booleans win.
P4	Tools never inherit	Tools and executable plugins require child declaration and host approval.

ID	Principle	Normative meaning
P5	Disabled means absent	Disabled skills and plugins do not enter prompts, tools, memory, cache keys, traces, evidence graphs, or critique.
P6	Persona is presentation	Persona cannot mint facts, permissions, confidence, credentials, trust, or authority.
P7	Evidence over narrative	Provenance records observable evidence, actions, constraints, and outcomes—not claimed private thought.
P8	Optimize job completion	Optimize successful, SLO-meeting work per time and cost, not raw token rate.
P9	Capability negotiation	Runtime behavior depends on verified capabilities, not vendor branding.
P10	Typed extensibility	Every extension has a manifest, typed interfaces, integrity data, permissions, and limits.
P11	Memory is governed data	Persistent memory has source, time, scope, trust, sensitivity, retention, and deletion semantics.
P12	Improvement proposes	The agent may generate candidates; only an external gate may promote them.
P13	Reproducible composition	Every run references immutable compose and capability locks.
P14	Fail closed at trust boundaries	Unknown code, drift, invalid memory, broken provenance, or unauthorized mutation aborts.
P15	Optional acceleration fails safely	Optional optimizer failure may use only a validated semantic-equivalent baseline.
P16	No floating production dependencies	Protocol, adapter, plugin, model, prompt, template, tokenizer, schema, and semconv versions are pinned.
P17	Independent validation	A self-generated candidate cannot promote on its own self-score.
P18	Rollback is mandatory	Every promoted mutable artifact has a previous signed version and tested rollback.
P19	Cache is scheduled	Cache has an owner, scope, budget, key discipline, correctness guard, and telemetry.
P20	Context is a lifecycle	Context is allocated, compacted, offloaded, re-grounded, and verified under policy.
P21	Capabilities are verified	An asserted capability is unavailable until a conformance fixture passes.
P22	Attribution is claim-level	Material output claims link to supporting evidence.

ID	Principle	Normative meaning
P23	Authority is by handle	Extensions receive narrow object-capability handles, not ambient authority.
P24	Memory is an attack surface	Persistent memory is treated as adversarially reachable.
P25	Every fixed failure becomes permanent	A resolved failure becomes a retained regression fixture.
P26	Verifier before optimizer	No improvement objective is admissible without an independent verifier.
P27	Corrigibility is unreachable by construction	Gates, telemetry, permissions, and termination controls are outside every agent-writable surface.
P28	Statistical honesty	Claims require adequate power, paired design, intervals, declared estimands, and pre-registration.
P29	Citation integrity	An unaudited or unverifiable citation cannot support release.
P30	Fail back or fail stop—never fail open	Optional optimizers fail back; mandatory controls stop execution.

4. Normative architecture

4.1 Nine planes



```
C --> O["Observability plane"]
E --> O
R --> O
T --> O
SAF --> O
CX --> O
O --> OT["OTel core + pinned GenAI semconv
+ casops.* stable aliases"]

O --> I["Improvement plane
attribution / candidate generation"]
I --> S["Sandbox + independent evaluation"]
S --> G{"Statistical release gates"}
G -->|pass| CAN["Shadow / canary"]
CAN --> H{"Independent human approval"}
H -->|"signed artifact"| F
H -->|reject| Q["Archive / quarantine"]

COR["Corrigibility plane
unreachable invariants"] -.->|"constrains"| E
COR -.->|"constrains"| I
COR -.->|"constrains"| O
COR -.->|"constrains"| R

H -. never automatic .-> PROD["Production activation"]
```

4.2 Control boundaries

Plane	May read	May write	Cannot do
Composer	Folder, parent metadata, manifests, conformance results	Generated locks	Execute unverified extension code
Execution	Composed envelope, approved memory and capabilities	Run artifacts, candidate memories	Modify its production definition
Cache/context	Scoped prefixes, entries, context state	Scoped cache and context artifacts	Cross tenant, subject, sensitivity, or approval boundaries
Memory	Approved observations and outcomes	Versioned records and tombstones	Grant tools, alter policy, or override newer human approval
Safety	Inbound content, external output, run state	Taint, blocks, incidents	Be bypassed by the agent, skill, plugin, or persona
Observability	Operational events and configured content	Append-only telemetry	Change run behavior
Improvement	Traces, outcomes, fixtures	Candidate artifacts	Promote, approve, or sign a candidate

Plane	May read	May write	Cannot do
Corrigibility	Host-held invariant reference	Attestation only	Be modified by any runtime plane
Human gate	Candidate and validation report	Approval record	Bypass immutable audit

4.3 Control-switch classes

Switch class	Applies to	Effect
optimizer_kill_switch	Cache tier, speculation, learned router, compaction strategy, adaptive refinement	Disable feature and return to validated baseline semantics
route_quarantine	Model, protocol, plugin, or capability route	Remove route from eligibility
containment_stop	Safety, corrigibility, mandatory audit, permission enforcement	Halt or reject work; never bypass the control
operator_shutdown	Entire agent	Cancel in-flight work within the configured deadline

5. Folder contract

5.1 Tree

```
agents/<pack.agent-id>/
  README.md
  SPEC.md
  agent_spec.json

  prompts/
  rubrics/
  sources/{PROVENANCE.json,MAPPING.md,excerpts/}
  docs/user_guide.md

  inheritance/{parents.json,resolved.json,conflicts.json}
  skills/{SKILL.md,bindings.json,integration.json,toggles.json}
  identity/{persona.json,background.json,DISCLOSURE.md}

  runtime/
    execution.json
    backends.json
    routing.json
    cache.json
    context.json
    compute_controller.json

  protocols/
```

```
compatibility.json
capability_assertions.json
conformance/
schemas/{agent_message.schema.json,event.schema.json}

observability/
  telemetry.json
  redaction.json
  slo.json
  decision_record.schema.json
  sampling.json
  evidence_graph.schema.json
  semconv.lock.json

plugins/
  registry.json
  lock.json
  manifests/
  isolation.json
  supply_chain/

memory/
  policy.json
  stores.json
  retention.json
  hierarchy.json
  consolidation.json
  security.json
  unlearning.json
  schemas/memory_record.schema.json
  migrations/

improvement/
  policy.json
  objectives.json
  verifiers.json
  ledger.json
  candidates/
  approvals/
  rollback/

safety/
  policy.json
  injection.json
  termination.json
  incidents/

corrigibility/
  invariants.json
  attestation.json

evals/
  benchmarks.json
  baselines.json
```

```
analysis_plan.json
fixtures/
regression/
reports/

generated/
compose.lock.json
capabilities.lock.json
benchmark-baseline.json
compatibility-matrix.lock.json
context-budget.lock.json
```

`corrigibility/invariants.json` is a logical folder contract path. At runtime it must be a host-owned, read-only mount or immutable reference outside every agent-writable capability.

5.2 Required files

Path	Requirement
README.md, SPEC.md, agent_spec.json	Always required
sources/PROVENANCE.json	Always required
inheritance/parents.json	Required; may be empty
inheritance/resolved.json	Generated after successful compose
skills/toggles.json	Required; may be empty
identity/DISCLOSURE.md	Required when identity mode is not grounded
runtime/execution.json	Required
runtime/backends.json	Required; may declare only local_deterministic
runtime/routing.json, runtime/cache.json	Required; disabled modes are valid
runtime/context.json	Required
runtime/compute_controller.json	Required; mode: fixed is valid
protocols/compatibility.json	Required
protocols/capability_assertions.json	Required
protocols/conformance/	Required and non-empty before production
observability/telemetry.json	Required
observability/redaction.json	Required
observability/sampling.json	Required
observability/semconv.lock.json	Generated; required before export
plugins/registry.json	Required; empty list is valid

Path	Requirement
plugins/isolation.json	Required when plugins are non-empty
memory/policy.json	Required; <code>mode: none</code> is valid
memory/hierarchy.json	Required when persistent memory is enabled
memory/security.json	Required when persistent memory is enabled
memory/unlearning.json	Required when persistent memory is enabled
improvement/policy.json	Required; defaults to disabled
improvement/verifiers.json	Required when improvement is enabled
safety/policy.json	Always required; no opt-out
safety/termination.json	Always required
corrigibility/invariants.json	Always required; host-owned and agent-unwritable
evals/benchmarks.json	Always required
evals/analysis_plan.json	Required before any performance or quality claim
evals/regression/	Required and union-monotonic
generated/*.lock.json	Generated only

5.3 Self-contained meaning

The folder must independently describe:

- mission and boundaries;
- expected inputs and outputs;
- inheritance;
- skills and identity;
- required and forbidden capabilities;
- runtime, cache, context, and compute policies;
- protocols and compatibility;
- telemetry and redaction;
- plugin and isolation policy;
- memory, trust, retention, security, and deletion;
- improvement scopes and verifiers;
- safety and termination;
- corrigibility invariants;
- validation and statistical analysis.

Parent folders, remote catalogs, endpoints, and host-held invariants are resolved only during compose.

6. Composition and inheritance

6.1 Parent legality and MRO

A parent is legal only if:

- 1. its `agent_id` resolves to a registered folder;
- 2. the folder contains `SPEC.md` and `agent_spec.json`;
- 3. it belongs to a compatible common-agent structure family;
- 4. it is not the child;
- 5. every requested inherited surface is legal;
- 6. its content hashes can be pinned.

Rules:

- maximum declared parents: 8;
- maximum depth: child plus three ancestor levels;
- child is first in MRO;
- parents sort by ascending priority, then ascending `agent_id`;
- traversal is depth-first;
- each parent appears once;
- diamonds collapse;
- cycles fail closed;
- missing or incompatible parents fail closed.

6.2 Merge rules

Surface	Rule
Child scalar content	Child wins
<code>does_not_own</code>	Union; more restrictive
Numeric budgets	Minimum
<code>network_access</code>	False wins
Production activation	False wins
<code>max_refinement_count</code>	Minimum
<code>allowed_tools</code>	Never inherited
<code>allowed_plugins</code>	Never inherited
Content lists	Stable-key union in MRO order
Critique edges	Union; self-edges removed
Safety policy	Tightening may inherit; relaxation does not
Prompt/rubric references	Child remains primary; parent references are labelled appendices
Parent SPEC fragments	Quoted with parent ID and hash, never silently represented as child-authored

6.3 Non-inherited surfaces

The following never inherit:

- credentials, secrets, API keys, and cookie jars;
- model endpoint credentials;
- provider authorization;
- protocol authorization;
- tools and executable plugin grants;
- plugin signatures and approvals;
- cache contents or cache-sharing permissions;
- verified capability status;
- tokenizer or template verification results;
- isolation-tier assignments or sandbox grants;
- SBOM and build-provenance approvals;
- actual memory records;
- tenant, subject, or user scope;
- memory tier residency;
- memory trust assignments;
- deletion or unlearning authority;
- telemetry content-capture approval;
- improvement approval or promotion rights;
- verifier approvals;
- canary allocation;
- production activation;
- named-person identity approval;
- safety-policy relaxations;
- termination-condition relaxations;
- corrigibility invariants;
- statistical plans, baselines, or held-out sets.

6.4 Legal inherited surfaces

Surface	Behavior
responsibility_fragments	Appended as support; child mission remains first
knowledge_sources	Stable-key union
quality_criteria	Union by dimension; child weight wins
prompt_refs	Labelled inherited appendices
rubric_refs	Labelled inherited appendices
skill_bindings	Union, then enable-AND
critique_edges	Union
architecture_pattern	Child wins; parents become inherited hints
persona_defaults	Used only if child identity files are absent
runtime_hints	Non-binding

Surface	Behavior
context_hints	Non-binding
protocol_preferences	Host and child policy win
plugin_requirements	Dependency declaration only
memory_schema_refs	Schema union only
evaluation_dimensions	Union by ID
verifier_refs	Definition union; approval does not inherit
regression_fixtures	Union-monotonic
safety_fixtures	Union-monotonic
failure_taxonomy	Union by code
observability_labels	Namespaced union
docs	Authoring hint only; not injected at runtime

Unknown surfaces fail closed.

6.5 Fixture monotonicity

FR-INH-301. Regression and safety fixtures are union-monotonic. A child may not remove an inherited fixture.

Removal requires:

- a signed host-level waiver;
- reason;
- expiry;
- impact assessment;
- compensating fixture or control;
- entry in `inheritance/conflicts.json`;
- human approval.

6.6 Compose lock

`generated/compose.lock.json` contains:

- child and parent hashes;
- MRO;
- prompt and rubric hashes;
- skill resolution;
- cache-policy hash;
- context-policy hash;
- compute-controller hash;
- verified capability matrix;
- tokenizer and chat-template digests;

- protocol and adapter versions;
- OTel `schema_url`;
- plugin digests;
- isolation-tier map;
- plugin SBOM digests;
- memory policy, hierarchy, security, and unlearning hashes;
- verifier-set hash;
- safety and termination hashes;
- corrigibility-invariant digest;
- analysis-plan hash;
- regression-fixture-set digest;
- `compose_hash`.

Any change creates a new `compose_hash`.

7. Performance execution plane

7.1 Runtime design

```

flowchart LR
    U["Task"] --> ADM["SLO admission control  
deadline / budget / risk"]
    ADM -->|"reject or queue"| SHED["Bounded response  
reason code"]
    ADM --> CC["Compute controller  
test-time budget"]
    CC --> MR["Model router"]
    MR --> PL["Planner"]
    PL --> DAG["Typed execution DAG"]

    DAG --> N1["Model node"]
    DAG --> N2["Memory node"]
    DAG --> N3["Tool node A"]
    DAG --> N4["Tool node B"]
    N1 --> J["Join"]
    N2 --> J
    N3 --> J
    N4 --> J

    J --> V["Deterministic validators"]
    V --> SG["Safety gate"]
    SG --> Q{"Pass?"}
    Q -->|"yes"| O["Output + evidence graph"]
    Q -->|"no; gain expected"| STOP{"Marginal gain > cost?"}
    STOP -->|"yes"| RF["Bounded refinement"]
    RF --> DAG
    STOP -->|"no"| X["Fail closed or escalate"]
    Q -->|"unsafe"| X
  
```

```

    CACHE[("Cache + context plane")] --- MR
    CACHE --- N1
    SCH["Deadline-aware scheduler
    goodput objective"] --- DAG

```

7.2 Execution IR

The execution IR is `casops.execution_dag.v2`.

Node kinds:

- `model;`
- `tool;`
- `plugin;`
- `memory_read;`
- `memory_write;`
- `peer_agent;`
- `validator;`
- `verifier;`
- `safety_check;`
- `compaction;`
- `branch;`
- `join;`
- `transform;`
- `speculative;`
- `human_gate.`

Every node declares:

- stable `node_id`;
- dependencies;
- typed input and output schemas;
- timeout;
- retry policy;
- side-effect class;
- idempotency;
- cacheability;
- required verified capabilities;
- resource budget;
- context-cost estimate;
- cache-affinity key;
- isolation tier where applicable;
- taint class;
- failure action;
- compensating action for speculative side effects.

7.3 Functional requirements

Base execution requirements

ID	Requirement
FR- PERF- 001	Explicit dependencies must be compiled before execution.
FR- PERF- 002	Independent read-only or idempotent nodes should run concurrently.
FR- PERF- 003	Unordered side-effecting nodes must not be parallelized.
FR- PERF- 004	The scheduler must optimize end-to-end critical-path time.
FR- PERF- 005	Every run propagates a deadline and cancellation token.
FR- PERF- 006	The router optimizes a declared quality-latency-cost objective.
FR- PERF- 007	Route decisions are reproducible from logged features or explicitly stochastic.
FR- PERF- 008	Structured output uses constrained decoding when verified and supported.
FR- PERF- 009	Prefix/KV caches are keyed by model, tokenizer, template, prompt, capability, tenant, subject, sensitivity, and approval scope.
FR- PERF- 010	Cache data must not cross agent, user, tenant, sensitivity, or approval boundaries.
FR- PERF- 011	Context reserves separate budgets for policy, task, memory, tools, evidence, and output.
FR- PERF- 012	Reflection is triggered by uncertainty, failure, risk, or expected utility—not an unconditional loop.

ID	Requirement
FR- PERF- 013	Speculative decoding must preserve the backend semantics it claims.
FR- PERF- 014	Optional optimizer failure may fall back only to a validated baseline.
FR- PERF- 015	Every fallback emits telemetry and a reason code.
FR- PERF- 016	Model, tool, memory, plugin, and peer concurrency have independent limits.
FR- PERF- 017	Acceptance jointly evaluates success, p50/p95 job time, and cost per successful task.

v3 requirements

ID	Requirement
FR- PERF- 101	SLO-aware admission control queues or sheds work rather than degrading all in-flight runs.
FR- PERF- 102	The scheduling objective is goodput, not throughput or tokens per second.
FR- PERF- 103	A compute controller allocates test-time budget from risk, deadline, value, and predicted difficulty and enforces a marginal-gain stopping rule.
FR- PERF- 104	Stopping decisions record estimated gain, cost, threshold, and rule version.
FR- PERF- 105	Speculative nodes declare a guard and cannot commit side effects before it passes. Abandoned speculation executes compensation.
FR- PERF- 106	Critical Path Efficiency is computed per run; sustained low CPE is treated as a scheduling defect.

ID	Requirement
FR- PERF- 107	Router records contain feature vector, candidates, scores, and selection-rule version.
FR- PERF- 108	Accelerator utilization is reported only when exposed by the backend; otherwise it is unavailable .
FR- PERF- 109	Every optional optimizer has an independent, fixture-tested kill switch. Mandatory safety, corrigibility, audit, and permission controls have no bypass switch; their failure invokes containment stop.
FR- PERF- 110	Acceptance jointly evaluates CPST, p50/p95/p99 where powered, task success, and goodput.

7.4 Metrics

Metric	Definition
CPST	Attributed model, tool, and infrastructure cost divided by successful tasks
goodput	Tasks that succeed and meet deadline per wall-second per accelerator
CPE	Ideal critical-path duration divided by actual wall time
CRR	Reused prefix tokens divided by total prompt tokens
TTF0	Time to first output token or artifact byte
job_completion_ms	Admission to artifact sealing
refinement_yield	Success-rate delta attributable to refinement divided by refinement cost share

7.5 Example compute controller

```
{
  "schema_version": "3.0",
  "agent_id": "video.showrunner",
  "mode": "adaptive",
  "budget_source": [
    "risk_class",
    "deadline",
    "predicted_difficulty",
    "task_value"
  ],
  "allocation": {
    "min_model_calls": 1,
    "max_model_calls": 6,
    "max_refinements": 3,
    "max_parallel_samples": 2
  }
}
```

```

    },
    "stopping_rule": {
      "type": "marginal_gain_threshold",
      "estimator": "validator_score_delta_ema",
      "min_expected_gain": 0.05,
      "cost_normalizer": "cpst",
      "hard_stop_on": [
        "deadline_80pct",
        "budget_90pct",
        "no_gain_2_consecutive"
      ]
    },
    "audit": {
      "log_gain_estimates": true,
      "log_threshold_version": true
    }
  }
}

```

7.6 Per-run fields

Every run records:

- admission_wait_ms;
- shed_reason;
- job_completion_ms;
- time_to_first_output_ms;
- critical_path_ms;
- queue_ms;
- model_ms;
- tool_wait_ms;
- memory_ms;
- validation_ms;
- token and cost totals;
- call and retry counts;
- cache tier hits;
- compute_budget_allocated;
- compute_budget_used;
- stopping_rule_decisions[];
- cpe;
- crr;
- context_tokens_by_segment;
- compaction_events[];
- committed and discarded speculation;
- kill_switch_engaged[];
- containment_stop_reason;
- goodput_contribution;
- cpst_contribution;
- outcome and failure code.

8. Cache and context-lifecycle plane

Evidence basis retained from the prior draft: agent-aware cache-management and agentic-serving work [D]. All such evidence is E3 until audited.

8.1 Rationale

Agent runs often reuse:

- long system and policy prefixes;
- tool schemas;
- prompt templates;
- memory excerpts;
- multi-turn context;
- peer-agent artifacts.

Cache can improve performance but creates high-risk failure modes:

- cross-tenant leakage;
- policy staleness;
- approval staleness;
- memory-deletion residue;
- tokenizer or template mismatch;
- approximate semantic false reuse.

Cache is therefore a scheduled, governed plane.

8.2 Cache tiers


Tier	Contents	Correctness guard
T0	Tokenized prefix or KV state	Exact key including policy, model, tokenizer, and template digests
T1	Rendered fragments	Source-record version, tenant, subject, trust, and sensitivity
T2	Pure/idempotent node results	Full typed input digest and adapter revision
T3	Approximate semantic reuse	Off by default; independent equivalence verifier and false-reuse gate

8.3 Cache requirements

ID	Requirement
FR-CACHE-001	Keys include model revision, tokenizer digest, chat-template digest, prompt/policy digest, capability scope, agent, tenant, subject, sensitivity, and approval epoch.

ID	Requirement
FR-CACHE-002	Entries cannot cross agent, user, tenant, sensitivity, or approval boundaries. Violation aborts and purges.
FR-CACHE-003	Policy, prompt, memory-version, template, tokenizer, capability, or approval changes invalidate dependencies before the next read.
FR-CACHE-004	T3 remains disabled unless an equivalence verifier exists and false reuse is at most 0.5% on the declared fixture.
FR-CACHE-005	Cache-enabled and cache-disabled execution must be correctness-equivalent within the pre-registered margin.
FR-CACHE-006	Every tier has a declared budget and eviction policy. Eviction must not create silent staleness.
FR-CACHE-007	Hits, misses, invalidations, evictions, and scope rejections emit reason-coded telemetry.
FR-CACHE-008	Failure falls back to the uncached validated path without changing permissions or semantics.
FR-CACHE-009	Memory deletion propagates to all cache tiers.

8.4 Context lifecycle



```
J --> K["Re-grounding checkpoint"]
K --> J
```

8.5 Context requirements

ID	Requirement
FR- CTX- 001	Context is segmented with independent budgets and actual usage recording.
FR- CTX- 002	Safety charter, corrigibility constraints, does_not_own , disclosure, output schema, and active deadline are pinned and non-compactable.
FR- CTX- 003	Compaction runs a preservation verifier. Failure escalates or stops; it never silently proceeds.
FR- CTX- 004	Compacted content is offloaded to a retrievable reference and not destroyed mid-run.
FR- CTX- 005	Oversized subtasks should use isolated, narrowly briefed sub-agents rather than unbounded parent context.
FR- CTX- 006	Long-horizon runs insert configured re-grounding checkpoints.
FR- CTX- 007	A context-rot fixture demonstrates non-inferiority to an oracle-short-context condition within the declared margin.

9. Compatibility and protocol plane

9.1 Canonical interfaces

- **ModelAdapter**
- **ToolAdapter**
- **PeerAdapter**
- **MemoryAdapter**
- **TelemetryAdapter**
- **PluginRuntime**
- **EventAdapter**
- **CacheAdapter**
- **VerifierAdapter**

- **SafetyAdapter**

9.2 Base compatibility requirements

ID	Requirement
FR-CMP-001	Protocol and adapter versions are pinned.
FR-CMP-002	Adapters publish machine-readable capability assertions.
FR-CMP-003	Unknown capabilities fail closed.
FR-CMP-004	Vendor fields remain under vendor_extensions .
FR-CMP-005	Vendor extensions cannot override host gates.
FR-CMP-006	Message payloads have JSON Schema or an approved binary schema.
FR-CMP-007	Incompatible major schema versions are rejected.
FR-CMP-008	Bridges preserve trace, deadline, identity, authorization, and taint.
FR-CMP-009	Discovery results are cached only within advertised validity.
FR-CMP-010	Protocol adapters pass conformance fixtures before production.

9.3 Asserted versus verified capabilities



ID	Requirement
FR-CMP-101	Every capability is VERIFIED , REFUTED , or ASSERTED_UNVERIFIED .
FR-CMP-102	Only VERIFIED capabilities may bind in production.
FR-CMP-103	Conformance reruns after endpoint, model, adapter, tokenizer, template, protocol, or material configuration changes.

ID	Requirement
FR-CMP-104	A previously verified capability that fails is quarantined as <code>CMP_CAPABILITY_DRIFT</code> .
FR-CMP-105	Tokenizer and chat-template digests are pinned and treated as semantic compatibility inputs.
FR-CMP-106	Structured output negotiates an explicit JSON-Schema subset profile. Unsupported constructs fail compose.
FR-CMP-107	Seed or determinism claims require repeat-run verification. Otherwise they are <code>best_effort</code> .

Capability vocabulary includes:

```
chat
text_generation
streaming
tool_calls
parallel_tool_calls
parallel_tool_calls_verified
structured_output
json_schema
json_schema_profile
embeddings
reranking
vision
audio_input
audio_output
logprobs
seed
token_count
token_count_exact
prefix_cache
prefix_cache_explicit
kv_cache_reuse_across_requests
speculative_decoding
batching
cancellation
cancellation_mid_stream
context_length_verified
tool_choice_forcing
batch_invariant_kernels
```

9.4 Telemetry compatibility

ID	Requirement
----	-------------

ID	Requirement
FR-CMP-108	The semantic-convention <code>schema_url</code> is pinned in <code>semconv.lock.json</code> and the compose lock.
FR-CMP-109	Every gate-bearing GenAI attribute is emitted under both the pinned external name and a stable <code>casops.*</code> alias. Gates bind to <code>casops.*</code> .
FR-CMP-110	A semconv version change raises <code>CMP_SEMCONV_VERSION</code> and requires conformance rerun and alias-map review.
FR-CMP-111	The alias map is committed, versioned, and cannot change silently.

This requirement does not depend on proving that an external convention is currently unstable; it protects CASOPS from any external schema change.

9.5 MCP tool protocol

The prior draft lists MCP revisions `2025-03-26`, `2025-06-18`, `2025-11-25`, and `2026-07-28` [D].

ID	Requirement
FR-CMP-112	Exact MCP revision and SDK digest are pinned.
FR-CMP-113	The host supports the pinned revision and at least one prior supported revision.
FR-CMP-114	Unknown major versions fail closed; unknown minor additions are ignored, not inferred.
FR-CMP-115	Extensions are explicitly allow-listed; unrecognized extensions are inert data.
FR-CMP-116	Discovery is not authorization; discovered but unapproved tools remain unreachable.
FR-CMP-117	Discovery caches expire at advertised validity and invalidate on revision change.

9.6 Peer protocol

A2A remains the preferred external peer adapter, normalized into the CASOPS envelope.

ID	Requirement
FR-CMP-118	Peer envelopes carry a taint class; external-peer content is data, not instruction.
FR-CMP-119	Bridges preserve trace, deadline, authorization, and taint.
FR-CMP-120	Authorization is explicit, non-transitive, and cannot be widened by a bridge.
FR-CMP-121	Multi-agent exchanges enforce hop, cost, time, and cycle guards.

9.7 Canonical peer envelope

```

{
  "schema_version": "3.0",
  "message_id": "msg_01",
  "conversation_id": "conv_01",
  "task_id": "task_01",
  "from_agent": "video.showrunner",
  "to_agent": "video.editor",
  "message_type": "handoff",
  "hop_count": 2,
  "max_hops": 6,
  "parts": [
    {
      "kind": "data",
      "schema": "video.cut_brief.v3",
      "content_ref": "artifact://cut-brief/123"
    }
  ],
  "taint": {
    "class": "external_peer",
    "instruction_authority": false
  },
  "traceparent": "00-4bf92f3577b34da6a3ce929d0e0e4736-00f067aa0ba902b7-01",
  "deadline": "2026-08-24T20:00:00Z",
  "budget_remaining": {
    "cost_units": 12.5,
    "wall_ms": 18000
  },
  "auth_scope": [
    "artifact:read:cut-brief"
  ],
  "provenance": {
    "compose_hash": "sha256:...",
    "artifact_hash": "sha256:..."
  }
}

```

Private reasoning, unrestricted memory, credentials, and undeclared tool handles must not be transmitted.

10. Observability and decision provenance

10.1 Model

```

flowchart TB
    RUN["Run root span"]
    RUN --> CMP["Compose span"]
    RUN --> ADM["Admission span"]
    RUN --> PLAN["Planning span"]
    RUN --> ROUTE["Routing span"]

```

```

    RUN --> CTX["Context spans"]
    RUN --> MODEL["Model spans"]
    RUN --> TOOL["Tool / plugin spans"]
    RUN --> MEM["Memory spans"]
    RUN --> PEER["Peer spans"]
    RUN --> SAFE["Safety spans"]
    RUN --> VAL["Validation spans"]
    RUN --> IMP["Improvement span"]

    MODEL --> DR["Decision records"]
    TOOL --> DR
    MEM --> DR
    SAFE --> DR
    VAL --> DR

    DR --> EG["Evidence graph
claim to evidence"]
    EG --> ES["Append-only hash-chained store"]
    DR --> ES
    ES --> SAMP["Tail sampling"]
    SAMP --> OTEL["OTel exporter
pinned schema + aliases"]
    ES --> REPLAY["Replay / counterfactual audit"]
    ES --> RCA["Failure classification"]

```

10.2 Decision provenance

Every material decision record contains:

- decision, trace, and span IDs;
- agent and compose hash;
- objective;
- observable input references or hashes;
- available actions;
- selected action;
- machine-readable rejection or selection codes;
- evidence and memory references;
- policy results;
- validator and verifier results;
- calibrated confidence where available;
- prompt, model, tokenizer, and template digests;
- timestamps and outcome.

Raw hidden chain-of-thought is neither required nor authoritative.

10.3 Internal reasoning monitor

ID	Requirement
----	-------------

ID	Requirement
FR-OBS-101	An internal-only reasoning-monitor channel may be enabled for automated safety monitoring.
FR-OBS-102	Its content cannot enter exports, artifacts, memory, peer messages, prompts, evidence graphs, or telemetry payloads.
FR-OBS-103	Retention defaults to at most 24 hours and is encrypted.
FR-OBS-104	Only structured verdicts—risk, category, confidence—may enter telemetry.
FR-OBS-105	Verdicts may block execution but cannot support factual output claims.

The agent itself cannot read this channel or influence its verdicts.

10.4 Claim-level evidence graph

```
{
  "schema_version": "3.0",
  "artifact_id": "artifact://cut-brief/123",
  "trace_id": "4bf92f3577b34da6a3ce929d0e0e4736",
  "claims": [
    {
      "claim_id": "c1",
      "span": {
        "start": 240,
        "end": 318
      },
      "text_hash": "sha256:...",
      "claim_type": "constraint",
      "support": [
        {
          "kind": "memory",
          "ref": "memory://semantic/17",
          "record_version": 3,
          "strength": "direct"
        },
        {
          "kind": "source",
          "ref": "source://director-spec/sha256:...",
          "strength": "direct"
        }
      ],
      "taint_inherited": "none",
      "verifier": {
        "name": "constraint_grounding_v2",
        "result": "pass",
        "score": 0.94
      }
    }
  ]
}
```

```
    }
  },
  {
    "claim_id": "c2",
    "span": {
      "start": 402,
      "end": 466
    },
    "claim_type": "inference",
    "support": [
      {
        "kind": "derived",
        "from": ["c1"],
        "strength": "inferred"
      }
    ],
    "verifier": {
      "name": "constraint_grounding_v2",
      "result": "flagged",
      "score": 0.41
    }
  }
],
"unsupported_claim_count": 0,
"flagged_claim_count": 1
}
```

ID	Requirement
FR-OBS-106	Claim-bearing artifacts emit an evidence graph.
FR-OBS-107	Claims resolve to source , versioned memory , tool_observation , derived , or unsupported .
FR-OBS-108	Tainted support propagates its taint to the claim.
FR-OBS-109	Unsupported-claim rate is reported per run and gated.
FR-OBS-110	A memory-write candidate derived from an unsupported claim is rejected.

10.5 Telemetry events

```
agent.run.started
agent.compose.completed
agent.admission.decided
agent.compute.budget_allocated
agent.compute.stop_decided
```

```
agent.route.selected
agent.plan.created
agent.node.started
agent.model.request
agent.model.response
agent.tool.request
agent.tool.response
agent.cache.hit
agent.cache.miss
agent.cache.invalidated
agent.cache.scope_rejected
agent.context.compaction_started
agent.context.compaction_verified
agent.context.regrounded
agent.capability.verified
agent.capability.drift_detected
agent.memory.query
agent.memory.result
agent.memory.page_in
agent.memory.page_out
agent.memory.write.proposed
agent.memory.write.committed
agent.memory.consolidation_completed
agent.memory.poison_suspected
agent.memory.unlearn_verified
agent.peer.message.sent
agent.peer.message.received
agent.safety.injection_detected
agent.safety.taint_propagated
agent.safety.blocked
agent.policy.decision
agent.validation.completed
agent.verifier.completed
agent.refinement.started
agent.evidence.graph_emitted
agent.improvement.candidate_created
agent.improvement.ledger_appended
agent.termination.guard_triggered
agent.corrigibility.attested
agent.corrigibility.violation_detected
agent.run.completed
agent.run.failed
```

Mandatory attributes include:

- identity, trace, span, correlation, and compose IDs;
- MRO;
- model and adapter revisions;
- tokenizer and template digests;
- semconv schema URL;
- capability-lock digest;
- prompt and skill hashes;

- plugin digests;
- tool, memory, and peer IDs;
- timings, tokens, cost, and outcome;
- cache tiers, CRR, CPE;
- context usage and compaction count;
- taint and trust tiers;
- verifier and safety results;
- sampling decision and reason;
- disclosure and expertise mode;
- corrigibility attestation ID.

Gate-bearing fields are emitted under `casops.*`.

10.6 Sampling and trace cost

ID	Requirement
FR-OBS-111	Agent-run sampling is tail-based.
FR-OBS-112	Failures, safety blocks, high-risk work, denials, memory writes, candidates, promotions, rollbacks, capability drift, and corrigibility events receive 100% retention.
FR-OBS-113	Successful low-risk runs may be sampled if the rate and reason are recorded.
FR-OBS-114	Trace-budget exhaustion degrades content capture first, then low-risk sampling. Mandatory categories cannot be dropped.
FR-OBS-115	Aggregate metrics use unsampled counters.

10.7 Content capture

Level	Behavior
<code>metadata_only</code>	Default: hashes, schemas, sizes, IDs
<code>redacted</code>	Approved fields after redaction
<code>encrypted_full</code>	Explicit approval, encryption, limited retention
<code>disabled</code>	Only if mandatory local action metadata remains

Prompts, outputs, memory content, tool arguments, and peer messages default to `metadata_only`.

10.8 Integrity, replay, and RCA

- Events are append-only and hash-chained.
- Deterministic runs record seed, model revision, tokenizer, template, environment, and fixtures.
- Without verified batch-invariant kernels, replay is observation-level rather than token-level.
- Exporter failure uses a bounded encrypted local spool.
- If exporter and local audit are both unavailable, high-risk execution containment-stops.
- Failed runs receive a versioned failure classification.
- **RCA@1** is measured on injected single-fault scenarios.
- Counterfactual replay may replace one route, memory, or tool observation.
- Counterfactual runs cannot write memory or publish production artifacts.

11. Extensible plugin architecture

11.1 Kinds

```
tool
skill_runtime
model_adapter
memory_backend
modality_handler
evaluator
protocol_adapter
validator
router
scheduler_policy
cache_adapter
verifier
safety_adapter
compaction_strategy
consolidation_job
```

Skills are declarative. Plugins are executable. A skill cannot install, sign, tier, grant, or authorize a plugin.

11.2 Isolation tiers

Tier	Mechanism	Threat model	Permitted side effects	Target overhead
I0	In-process	Trusted first-party only	Read-only	Negligible
I1	Capability-based WASM	Buggy or semi-trusted	Pure transform, read-only	≤1 ms median; ≤3% p95
I2	Separate process, namespace/seccomp, no ambient network	Semi-trusted third party	Scoped read/write	≤5% p95

Tier	Mechanism	Threat model	Permitted side effects	Target overhead
I3	MicroVM, no host FS, allow-listed egress proxy	Untrusted or adversarial	External effect	≤15% p95
ID	Requirement			
FR-PLG-101	Every plugin is assigned a tier.			
FR-PLG-102	Tier meets or exceeds side-effect and provenance minimums.			
FR-PLG-103	Third-party or unsigned-origin code cannot run below I2; network-capable code runs at I3.			
FR-PLG-104	Downgrade requires a signed, expiring host waiver.			

11.3 Object-capability authority

ID	Requirement
FR-PLG-105	Plugins receive narrow capability handles, never ambient credentials or general clients.
FR-PLG-106	Handles are unforgeable, revocable, expiring, and revoked at node completion.
FR-PLG-107	Plugins cannot enumerate capabilities they were not granted.
FR-PLG-108	Handle delegation is prohibited unless explicitly declared and approved.

11.4 Supply chain

ID	Requirement
FR-PLG-109	Every production plugin supplies an SBOM.
FR-PLG-110	Every production plugin supplies verifiable build provenance.
FR-PLG-111	Components pass the configured vulnerability threshold.
FR-PLG-112	Signatures verify against approved keys and digests match.
FR-PLG-113	Builds should be reproducible; exceptions require justification.

11.5 Base and lifecycle requirements

ID	Requirement
FR-PLG-001	Installation requires zero composer-core source changes.
FR-PLG-002	Every interface is typed.
FR-PLG-003	Every plugin declares permissions and limits.

ID	Requirement
FR-PLG-004	Undeclared network, file, memory, tool, or peer access is blocked.
FR-PLG-005	Plugin code cannot execute during manifest validation.
FR-PLG-006	Unknown or unsigned production plugins fail closed.
FR-PLG-007	Dependency cycles fail closed.
FR-PLG-008	Hot reload requires quiescence or a stateless declaration.
FR-PLG-009	Output remains untrusted until schema and policy validation pass.
FR-PLG-010	Tool plugins are evaluated on function-calling and stateful interaction fixtures [C].
FR-PLG-114	ABI versions use semantic versioning and contract tests.
FR-PLG-115	Deprecated interfaces have a declared support window, defaulting to at least two minor host releases.
FR-PLG-116	Hot swap drains in-flight calls and shadow-validates the replacement.
FR-PLG-117	Manifest validation never executes plugin code.
FR-PLG-118	Plugin output is taint-labelled until validated.

Lifecycle:

- 1. discover;
- 2. schema-validate;
- 3. verify digest/signature;
- 4. verify SBOM, provenance, and scan;
- 5. resolve dependencies;
- 6. verify compatibility and ABI;
- 7. assign isolation;
- 8. evaluate permissions;
- 9. instantiate without ambient authority;
- 10. health-check;
- 11. register typed capabilities;
- 12. lock and emit event;
- 13. quiesce before update or removal.

11.6 Manifest

```
{
  "schema_version": "3.0",
  "plugin_id": "video.frame-inspector",
  "version": "3.1.0",
  "kind": "modality_handler",
  "entrypoint": {
    "runtime": "wasm",
    "path": "plugins/frame-inspector/main.wasm"
  },
  "isolation": {
    "tier": "I1",
    "threat_model_ref": "docs/threat-model.md#i1"
  },
  "abi": {
    "interface_version": "2.1",
    "contract_tests": "conformance/frame-inspector/"
  },
  "interfaces": [
    {
      "name": "inspect_frame",
      "input_schema": "schemas/frame-request.json",
      "output_schema": "schemas/frame-result.json",
      "deterministic": true
    }
  ],
  "capabilities_required": [
    {
      "kind": "artifact_read",
      "scope": "artifact://frames/*",
      "handle": "narrow"
    },
    {
      "kind": "memory_read",
      "scope": "resource",
      "handle": "narrow"
    }
  ],
  "permissions": {
    "network": false,
    "filesystem_write": [],
    "tools": [],
    "memory_write": []
  },
  "side_effect_class": "read_only",
  "resource_limits": {
    "cpu_ms": 5000,
    "memory_mb": 512,
    "output_bytes": 1048576,
    "wall_ms": 8000
  },
}
```

```
"compatibility": {
  "agent_structure": ">=3.0 <4.0",
  "host_api": ">=5.0 <6.0"
},
"supply_chain": {
  "sbom_ref": "plugins/supply_chain/frame-inspector.sbom.json",
  "provenance_ref": "plugins/supply_chain/frame-inspector.provenance.json",
  "scan_result_ref": "plugins/supply_chain/frame-inspector.scan.json"
},
"integrity": {
  "sha256": "sha256:...",
  "signature_ref": "approval://plugin-signature/123"
}
}
```

12. Long-term memory architecture

12.1 Stores

Store	Purpose	Default duration
Working	Current task state	Run
Episodic	Attempts, outcomes, reflections	Configured
Semantic	Verified facts and relationships	Configured
Procedural	Reusable methods and workflows	Configured
Resource	Documents and multimodal artifacts	Configured
Profile/core	Explicit preferences and durable identity facts	Until changed or deleted
Evidence vault	Immutable sources, approvals, validation evidence	Policy-controlled

This is an engineering taxonomy, not a biological-equivalence claim.

12.2 Paged hierarchy

Tier	Residency	Target	Contents
H0	In context	Immediate	Pinned invariants, core profile, active state
H1	Indexed warm store	p95 ≤150 ms	Recent episodic and active semantic/procedural
H2	Cold archive	p95 ≤2 s	Historical, superseded, large resources
H3	Frozen evidence vault	Policy-defined	Sources, approvals, validation evidence
ID	Requirement		
FR-MEM-101	Every tier has a residency budget.		

ID	Requirement
FR-MEM-102	Page-in/out emits trigger, token cost, latency, and tier.
FR-MEM-103	H0 page-out preserves a retrievable reference and task-critical state.
FR-MEM-104	Pinned invariants are non-evictable from H0.
FR-MEM-105	Memory token cost is attributed by tier.

12.3 Record

```
{
  "schema_version": "3.0",
  "memory_id": "mem_01",
  "agent_id": "video.showrunner",
  "tenant_scope": "project:series-a",
  "subject_scope": "user:hashed-id",
  "memory_type": "semantic",
  "tier": "H1",
  "status": "verified",
  "content_ref": "encrypted://memory/mem_01",
  "summary": "Approved exterior-shoot curfew is 22:00 local time.",
  "entities": [
    "location:lot-a",
    "policy:curfew"
  ],
  "relations": [
    {
      "subject": "location:lot-a",
      "predicate": "has_curfew",
      "object": "22:00"
    }
  ],
  "valid_time": {
    "from": "2026-08-01T00:00:00Z",
    "to": null
  },
  "transaction_time": "2026-08-24T12:00:00Z",
  "source_refs": [
    "artifact://approval/curfew-2026"
  ],
  "provenance_chain": [
    "human_approval:appr_77"
  ],
  "trust_tier": "T0_human_verified",
  "taint": {
    "class": "none",
    "instruction_authority": false
  },
  "confidence": {
    "value": 1.0,
    "basis": "human_approval"
  }
}
```

```
  },
  "sensitivity": "internal",
  "retention_class": "project_lifetime",
  "supersedes": [],
  "conflicts_with": [],
  "access_stats": {
    "reads": 41,
    "last_read": "2026-08-23T09:12:00Z",
    "contributed_to_success": 38
  },
  "quality": {
    "utility_score": 0.92,
    "staleness_risk": "low"
  },
  "created_by_trace": "trace_01"
}
```

12.4 Trust tiers

Tier	Origin	Treatment
T0_human_verified	Human approval or signed source	Authoritative constraint
T1_validator_verified	Deterministic validation	Usable with source
T2_source_grounded	Provenanced source	Usable with displayed provenance
T3_agent_inferred	Unverified inference	Advisory only
T4_quarantined	Untrusted or suspect	Not retrievable into prompts

FR-MEM-106. T3 cannot appear as factual evidence or override T0/T1.

12.5 Lifecycle

```
flowchart LR
  OBS["Observation / outcome"] --> TAINT["Taint + trust classify"]
  TAINT --> CAND["Candidate extraction"]
  CAND --> RED["Redact + sensitivity classify"]
  RED --> PROV["Attach provenance + time"]
  PROV --> POIS["Poisoning / injection screen"]
  POIS --> DEDUP["Deduplicate + conflict check"]
  DEDUP --> Q{"Verified?"}
  Q -->|no| QUAR["Quarantine T4"]
  Q -->|yes| STORE["Typed store + tier"]
  STORE --> LINK["Entity / relation / temporal graph"]
  LINK --> CONS["Offline consolidation"]

  QUERY["Task query"] --> QP["Query planner"]
  QP --> HYB["Lexical + dense + graph + temporal"]
  HYB --> RERANK["Rerank + diversify"]
  RERANK --> TRUST["Trust filter"]
```

```
TRUST --> VERIFY["Source + conflict verification"]
VERIFY --> INJECT["Budgeted page-in"]

STORE --> RET["Retention / decay / deletion"]
RET --> TOMB["Tombstone propagation"]
TOMB --> UNL["Unlearning verification"]
```

12.6 Retrieval

Retrieval supports:

- 1. metadata, agent, tenant, and subject filtering;
- 2. lexical retrieval;
- 3. dense retrieval;
- 4. optional graph traversal;
- 5. temporal filtering and query expansion;
- 6. deterministic or learned reranking;
- 7. diversity control;
- 8. token budgeting;
- 9. source verification;
- 10. conflict-aware abstention.

ID	Requirement
FR-MEM-107	Trust filtering occurs before context injection.
FR-MEM-108	Configured graph stores support validity-aware temporal traversal [K].
FR-MEM-109	Retrieval returns the latest valid version and material unresolved conflicts.
FR-MEM-110	Irreconcilable conflict or insufficient coverage produces abstention.
FR-MEM-111	Retrieval reports token cost and utility attribution.

12.7 Offline consolidation

ID	Requirement
FR-MEM-112	Consolidation runs as scheduled offline work, not inline in latency-bound serving.
FR-MEM-113	Consolidation output is a candidate and cannot exceed the trust of its lowest-tier input.
FR-MEM-114	Superseded originals remain until retention policy expires them.
FR-MEM-115	Consolidation has independent capacity and cannot consume serving reservations.

12.8 Write policy

Allowed writes include:

- explicit preferences;
- sourced verified facts;

- successful outcomes;
- approved procedures;
- validator-confirmed lessons;
- human corrections;
- versioned task state.

The system must not directly commit:

- raw chain-of-thought;
- unsupported persona claims;
- unsupported evidence-graph claims;
- untrusted tool or peer instructions;
- content without provenance;
- secrets outside approved vaults;
- inferred sensitive traits;
- cross-tenant data;
- transient errors as durable rules;
- T3 inference as verified fact;
- tainted content without independent verification;
- consolidation output above input trust.

12.9 Conflict, deletion, and unlearning

- Changes create new versions.
- Older records are superseded, not silently overwritten.
- Valid and transaction time remain distinct.
- Contradictions persist until resolved or expired.
- The agent cannot silently choose the convenient conflict.

ID	Requirement
FR-MEM-116	Tombstones propagate to records, indexes, caches, summaries, embeddings, graph edges, consolidation output, and flagged derived artifacts.
FR-MEM-117	Post-deletion probes verify absence through lexical, dense, graph, and cache paths.
FR-MEM-118	Completion and latency are audited against the retention SLA.
FR-MEM-119	Legal-hold evidence is excluded from decay/deletion and the exclusion is auditable.
FR-MEM-120	If memory-derived data entered model training, deletion records the weight-level limitation and flags retraining review.

12.10 Metrics

Metric	Definition
TCE	Memory tokens injected divided by correct answers attributable to memory
unsupported_memory_answer_rate	Memory-using answers without valid provenance divided by memory-using answers
staleness_rate	Answers using a superseded record when a current record existed
MPR	1 minus successful poisoning attacks divided by attempts
DCR	Verified-complete deletions divided by deletion requests
page_in_cost	Tokens and latency attributable to page-in

12.11 Evaluation

Required categories:

- extraction;
- multi-session reasoning;
- temporal reasoning;
- update;
- abstention;
- long-form and multimodal memory;
- retrieval;
- test-time learning;
- long-range understanding;
- selective forgetting;
- memory-to-action;
- poisoning resistance;
- deletion and unlearning verification;
- staleness after update.

A public benchmark score alone cannot satisfy the memory gate. Contamination checks and domain golden tasks are mandatory.

13. Autonomous self-improvement

13.1 Levels and writable locations

Level	Scope	Default	Writable location
L0	Disabled	Production default	Nothing

Level	Scope	Default	Writable location
L1	Per-run retry, reflection, bounded search	Budgeted	Run-local state
L2	Candidate memory, prompt, context, rubric, router	Propose-only	improvement/candidates/
L3	Candidate workflow or plugin	Sandbox only	Sandbox and candidates
L4	Model adapter or LoRA-style weights	Separate trainer	Trainer artifacts
L5	Core source or self-rewriting architecture	Research-only	Isolated research environment

Never writable:

- corrigibility;
- safety policy;
- termination policy;
- permissions;
- tools and plugin grants;
- credentials;
- telemetry mandatory-retention configuration;
- redaction policy;
- gate thresholds;
- held-out sets;
- analysis plans;
- production activation;
- isolation-tier downgrades;
- regression-fixture removal.

13.2 Loop

```
flowchart TB
    OUT["Task outcome"] --> CLS["Failure classification"]
    CLS --> ATTR["Specific cause"]
    ATTR --> VER{"Independent verifier exists?"}
    VER -->|no| BLOCK["Reject objective"]
    VER -->|yes| OBJ["Select measurable objective"]
    OBJ --> GEN["Generate candidate"]
    GEN --> SB["Isolated sandbox"]

    SB --> F1["Functional"]
    SB --> F2["Held-out quality"]
    SB --> F3["Safety"]
    SB --> F4["Latency / cost / goodput"]
    SB --> F5["Compatibility"]
    SB --> F6["Observability"]
    SB --> F7["Full regression suite"]
    SB --> F8["Reward-hacking detectors"]
```

```

F1 --> STAT{"Powered gates pass?"}
F2 --> STAT
F3 --> STAT
F4 --> STAT
F5 --> STAT
F6 --> STAT
F7 --> STAT
F8 --> STAT

STAT -->|no| ARC["Archive + ledger"]
STAT -->|yes| SHAD["Shadow then canary"]
SHAD --> HUM{"Independent human approval"}
HUM -->|approve| SIGN["Sign + version + ledger"]
SIGN --> DEP["Deploy"]
DEP --> MON["Monitor"]
MON -->|regression| RB["Automatic rollback"]
HUM -->|reject| ARC
ARC --> RAT["Failure becomes fixture"]
RB --> RAT

```

13.3 Attribution

"Task failed" is insufficient. Cause codes include:

- route failure;
- missing or incorrect memory;
- retrieval granularity;
- malformed tool call;
- plugin defect;
- workflow dependency;
- inadequate validation;
- prompt ambiguity;
- protocol incompatibility;
- budget exhaustion;
- context overflow;
- compaction loss;
- cache staleness;
- capability drift;
- injection compromise;
- termination trip;
- verifier gap.

13.4 Candidate types and generators

Candidate types:

```

memory_correction
prompt_patch
context_playbook

```

```
rubric_patch
router_update
workflow_patch
plugin_patch
protocol_mapping
model_adapter
evaluation_fixture
```

Generator	Scope	Evidence	Default
Reflective prompt/context evolution	Prompt and context candidates	E2/E3 [K]	Preferred for auditability
Workflow search	Workflow candidate	E2 [C]	Gated
Router optimization	Router update	E1/E2 [C]	Bounded and shadowed
Memory correction	Memory candidate	E2	L2 with provenance
Trajectory trainer	Model adapter	E3 [C]	L4, out of process
Fixture synthesis	Evaluation fixture	E1	Encouraged; no held-out access

The preference for reflective evolution is operational: readable diffs, reversibility, lower infrastructure burden, and direct human review. Any empirical sample-efficiency claim remains citation-gated.

13.5 Candidate requirements

Every candidate contains:

- candidate ID and parent version;
- mutation scope;
- generating traces and failure codes;
- exact diff or artifact;
- expected benefit;
- risk and blast-radius assessment;
- training, evaluation, and holdout hashes;
- model, tool, adapter, template, and tokenizer versions;
- benchmark, safety, cost, and latency results;
- verifier identity and independence attestation;
- reward-hacking detector results;
- full regression results;
- statistical report;
- corrigibility attestation;
- rollback artifact;
- signature state.

13.6 Verifier and reward integrity

ID	Requirement
FR-IMP-101	Every objective names a verifier.
FR-IMP-102	Model judges are independent by model family or independently calibrated prompt lineage.
FR-IMP-103	Detectors cover verbosity, judge mimicry, memorization, refusal gaming, format exploitation, and metric/golden-task divergence.
FR-IMP-104	Golden-task degradation rejects the candidate regardless of target-metric gain.
FR-IMP-105	Held-out sets are cryptographically isolated.
FR-IMP-106	Fixture rotation is host-controlled.

13.7 Failure-to-fixture ratchet

ID	Requirement
FR-IMP-107	Every confirmed attributable failure creates a regression fixture before fix promotion.
FR-IMP-108	Fixtures are permanent and union-monotonic absent signed waiver.
FR-IMP-109	Every promotion passes the full suite; "known flaky" is not an exemption.
FR-IMP-110	Suite size and pass rate appear in every validation report.

13.8 Learning separation

Trajectories may be exported to a trainer. Gradient updates must not execute in the serving process.

Online serving updates may include:

- bounded router statistics;
- cache TTL estimates;
- quarantined reflections;
- memory indexes;
- non-executable task statistics;
- propose-only context playbooks.

They must not include:

- base-weight mutation;
- unsigned adapter promotion;
- plugin replacement;
- core source mutation;
- permission, tool, network, gate, telemetry, safety, or termination changes.

13.9 Core self-modification

L5 requires:

- isolated environment;
 - no production credentials;
 - disabled or tightly simulated outbound network;
 - approved writable repositories only;
 - separated evaluation sets;
 - human approval for every promotion;
 - signed rollback.
- It is never a standard production capability.

13.10 Improvement ledger

FR-IMP-111. Candidate creation, evaluation, statistics, approval, rejection, promotion, monitoring, and rollback append to an immutable hash-chained ledger. The agent cannot rewrite promotion-boundary entries.

14. Safety and adversarial-robustness plane

14.1 Threat model

Threat	Vector	Primary controls
Direct injection	User input	Instruction/data separation, taint
Indirect injection	Documents, tools, web, memory, peers	Taint propagation, no instruction authority
Memory poisoning	Persisted malicious content	Trust tiers, quarantine, MPR gate
Tool or agent hijacking	Redirected external actions	Allow-lists, confirmation, hijack tests
Exfiltration	Induced secret or PII output	Egress allow-list, outbound scanning
Supply chain	Malicious plugin/dependency	SBOM, provenance, signature, isolation
Excessive agency	Loops, spend, side effects	Hard caps, termination, compensation
Multi-agent cascade	Peer amplification	Hop caps, cycle detection, shared budget
Control bypass	Candidate or plugin alters safeguards	Corrigibility and negative fixtures

14.2 Taint propagation

ID	Requirement
FR-SAF-001	Every item has a taint class: operator, user, tool, untrusted retrieval, external peer, or memory trust tier.
FR-SAF-002	Taint propagates through transforms, summaries, compaction, and consolidation.
FR-SAF-003	Tainted content has <code>instruction_authority:false</code> .
FR-SAF-004	Taint-derived external effects require allow-list match, deterministic validation, or human confirmation.
FR-SAF-005	Outbound content passes secret and PII scanning.
FR-SAF-006	Safety cannot be disabled or bypassed by agent, skill, plugin, or persona.

14.3 Termination and excessive-agency guards

ID	Requirement
FR-SAF-007	Every run has hard limits for time, cost, model calls, tool calls, peer hops, refinements, and plan expansions.
FR-SAF-008	Progress-free loops halt after a configured non-improvement window.
FR-SAF-009	Peer cycles halt.
FR-SAF-010	Guard trips return explicit bounded failure, never silent truncated success.
FR-SAF-011	Near a cap, partial result with disclosure is preferred to unsafe completion.

14.4 Release gates

Gate	Threshold
Indirect injection	Observed attack success $\leq 2\%$, with no privilege escalation or secret exfiltration
Memory poisoning	MPR $\geq 95\%$; no poisoned content reaches T0/T1
Tool hijack	No unapproved external-effect action
Exfiltration	No secret or PII egress
Termination	All loop, hop, cycle, and budget fixtures halt correctly

Gate	Threshold
Taint integrity	All laundering fixtures blocked
Multi-agent cascade	No cross-agent privilege or taint escalation

Zero-tolerance categories are literal operational gates. Reports must also provide exact binomial confidence bounds; zero observed events is not represented as proof of zero population risk.

14.5 Incident handling

FR-SAF-012. Safety blocks, suspected injection, exfiltration attempts, and guard trips create incident records containing:

- trace;
- taint chain;
- attempted action;
- detection path;
- disposition;
- operator notification;
- resulting regression fixture.

15. Corrigibility plane

15.1 Invariants

ID	Invariant
INV-01	The agent cannot modify permissions, tools, or plugin grants.
INV-02	It cannot modify safety or termination policy.
INV-03	It cannot modify mandatory telemetry retention or redaction policy.
INV-04	It cannot modify gate thresholds, held-out sets, or analysis plans.
INV-05	It cannot request production activation or grant network access.
INV-06	It cannot approve, sign, or promote candidates.
INV-07	It cannot delete or rewrite audit, ledger, or incident records.
INV-08	It cannot disable, degrade, or bypass safety.
INV-09	It cannot remove regression or safety fixtures.
INV-10	It cannot suppress, delay, or reorder shutdown, cancellation, or deadline signals.
INV-11	It cannot read the reasoning-monitor channel or influence verdicts.
INV-12	It cannot lower plugin isolation or forge capability handles.

15.2 Enforcement

ID	Requirement
FR-COR-001	Enforcement uses separate ownership, storage, and capability absence—not policy checks alone.
FR-COR-002	Every compose attests the invariant digest against a host-held reference.
FR-COR-003	Mismatch invokes immediate containment stop and operator alert. No degraded mode exists.
FR-COR-004	Shutdown and cancellation are honored at every node boundary and enforceably terminate plugin invocations.
FR-COR-005	Candidates touching invariant surfaces are rejected at generation time and alerted.
FR-COR-006	INV-01 through INV-12 each have a negative fixture. Untested invariants are assumed broken.

Mandatory controls have no bypass kill switch. Their only emergency switch is containment stop.

16. Skills, identity, and persona isolation

16.1 Skill resolution

```
resolved_enabled =
    declared_or_inherited
AND author_enabled
AND inherited_enabled
AND operator_toggle
AND host_permission
```

ID	Requirement
FR-SKL-001	Bindings include <code>skill_id</code> , source, path, and enabled state.
FR-SKL-002	Missing <code>enabled</code> defaults to true.
FR-SKL-003	Operator toggles record reason, actor, and time.
FR-SKL-004	A toggle cannot enable an undeclared or uninherited skill.
FR-SKL-005	Enabling cannot add tools outside child and host grants.
FR-SKL-006	Disabled skills do not enter prompts, tools, memory, cache, evidence, or critique.
FR-SKL-007	Artifacts record loaded and disabled skills.

ID	Requirement
FR-SKL-008	Duplicate IDs collapse and enabled state is AND-combined.
FR-SKL-009	Unknown skill IDs fail closed.
FR-SKL-010	Enabled skills with missing required files fail closed.

A skill may reference an approved plugin but cannot install, sign, tier, or authorize it.

16.2 Identity modes

- grounded;
- persona_overlay;
- mixed.

ID	Requirement
FR-IDN-001	Missing identity files imply grounded.
FR-IDN-002	Persona may define voice, register, temperament, values, taboos, and languages.
FR-IDN-003	Background may define title, domain, specialties, fictional experience range, methods, education style, and credit style.
FR-IDN-004	Background is fictional unless a provenance-linked human approval authorizes a real person.
FR-IDN-005	Named-person personas require an approval ID and digest.
FR-IDN-006	Persona cannot alter tools, network, budgets, permissions, safety, or activation.
FR-IDN-007	Every non-grounded artifact displays disclosure and disclosure_id.
FR-IDN-008	Ungrounded craft claims are persona_claim and cannot support factual gates.
FR-IDN-009	Persona cannot claim real professional licenses or legal standing.
FR-IDN-010	Child identity overrides parent persona defaults.
FR-IDN-011	User psychological profiles are not copied into agent identity.
FR-IDN-012	Listed languages constrain supported output locales but do not invent support.

Persona additionally cannot affect:

- route-quality labels;
- capability verification;
- cache scope;
- memory trust;
- taint;
- safety verdicts;
- termination;
- evidence support;
- verifier selection;
- corrigibility.

17. Compose and runtime algorithm

17.1 Compose

1. Validate folder and JSON Schemas.
2. Attest corrigibility invariants before executable resolution.
3. Validate identity and disclosure.
4. Resolve MRO and hashes.
5. Enforce fixture monotonicity.
6. Merge content and tightening-only safety fields.
7. Resolve skills.
8. Discover plugin manifests without execution.
9. Verify signatures, digests, SBOMs, provenance, scans, ABI, permissions, dependencies, and isolation.
10. Resolve model, cache, protocol, verifier, and safety adapters.
11. Run capability conformance and write the verified matrix.
12. Pin tokenizer, template, protocol, adapter, and semconv revisions.
13. Bind memory stores, hierarchy, security, retention, and unlearning.
14. Bind telemetry, sampling, redaction, and local spool.
15. Bind cache and context budgets.
16. Validate verifier independence.
17. Compile the execution DAG and validate side-effect safety.
18. Apply tools, network, tenant, budget, and production gates.
19. Generate locks and `compose_hash`.
20. Run preview, safety, and negative-invariant fixtures.
21. Permit execution only if mandatory checks pass.

17.2 Run

1. Start the root trace.
2. Admit, queue, or shed based on risk, deadline, budget, and capacity.
3. Allocate test-time compute.
4. Query memory, trust-filter, and page in under budget.
5. Select and record the route.

6. Build or load the DAG.
7. Execute safe nodes concurrently.
8. Validate and taint-label all external results.
9. Compact only with preservation verification.
10. Re-ground at checkpoints.
11. Refine only while expected gain exceeds cost.
12. Run the safety gate.
13. Produce output, evidence graph, provenance, and disclosure.
14. Commit only verified memory writes.
15. Record metrics and failure classification.
16. Optionally create, but never promote, candidates.
17. Create regression fixtures for confirmed failures.
18. Close trace and seal the artifact.

17.3 Prompt envelope

1. Host safety charter — pinned
2. Corrigibility and permission constraints — pinned
3. Protocol constraints
4. Disclosure — pinned
5. Persona voice
6. Child mission and `does_not_own` — pinned
7. Inherited support fragments
8. Taint-labelled, trust-filtered memory
9. Enabled skills
10. Child primary prompt
11. Labelled inherited prompts
12. Tool and plugin schemas
13. Output schema — pinned
14. Rubric and validators
15. Deadline and remaining budget — pinned

18. Data models

18.1 `agent_spec.json`

```
{
  "schema_version": "3.0",
  "structure_id": "casops.common_agent.v3",
  "agent_id": "video.showrunner",
  "status": "registered",
  "role": "ShowrunnerAgent (VA Domain Pack)",
  "allowed_tools": [],
  "allowed_plugins": [],
  "model_policy": {
    "provider": "local_deterministic",
```

```

    "model_id": "local-video-config",
    "network_access": false,
    "routing_allowed": true
  },
  "budget_policy": {
    "max_input_tokens": 4096,
    "max_output_tokens": 1536,
    "max_model_calls": 4,
    "max_tool_requests": 0,
    "max_job_ms": 45000,
    "max_cost_units": 8.0,
    "max_peer_hops": 4
  },
  "prompt_reference": "video.prompt.showrunner.v3",
  "rubric_reference": "video.rubric.showrunner.v3",
  "critique_edges": {
    "inputs": ["video.critic"],
    "outputs": ["video.judge"]
  },
  "max_refinement_count": 3,
  "production_activation_requested": false,
  "does_not_own": [
    "Credentials",
    "Silent production activation",
    "Another agent's exclusive craft output without handoff",
    "Automatic promotion of self-generated artifacts",
    "Modification of safety, telemetry, gates, permissions, or corrigibility",
    "Self-granting tools, plugins, network, or isolation downgrades"
  ],
  "inheritance_ref": "inheritance/parents.json",
  "identity_ref": "identity/",
  "skills_ref": "skills/bindings.json",
  "toggles_ref": "skills/toggles.json",
  "runtime_ref": "runtime/execution.json",
  "context_ref": "runtime/context.json",
  "compute_controller_ref": "runtime/compute_controller.json",
  "backends_ref": "runtime/backends.json",
  "cache_ref": "runtime/cache.json",
  "protocols_ref": "protocols/compatibility.json",
  "capability_assertions_ref": "protocols/capability_assertions.json",
  "observability_ref": "observability/telemetry.json",
  "sampling_ref": "observability/sampling.json",
  "plugins_ref": "plugins/registry.json",
  "isolation_ref": "plugins/isolation.json",
  "memory_ref": "memory/policy.json",
  "memory_hierarchy_ref": "memory/hierarchy.json",
  "memory_security_ref": "memory/security.json",
  "improvement_ref": "improvement/policy.json",
  "verifiers_ref": "improvement/verifiers.json",
  "safety_ref": "safety/policy.json",
  "termination_ref": "safety/termination.json",
  "corrigibility_ref": "corrigibility/invariants.json",
  "evals_ref": "evals/benchmarks.json",

```

```
"analysis_plan_ref": "evals/analysis_plan.json"
}
```

18.2 Artifact metadata

Every output includes:

- agent and structure IDs;
- compose and capability locks;
- MRO and parent hashes;
- skills and plugins;
- model, route, tokenizer, and template;
- semconv schema;
- plan and trace IDs;
- memory reads/writes and trust tiers;
- cache tier hits, CRR, and CPE;
- context use and compaction;
- compute budget and stopping rule;
- evidence-graph ID;
- unsupported-claim rate;
- taint and safety verdicts;
- termination state;
- verifier results;
- regression-suite digest;
- corrigibility attestation;
- goodput and CPST contribution;
- expertise mode and disclosure.

18.3 Improvement policy

```
{
  "schema_version": "3.0",
  "agent_id": "video.showrunner",
  "mode": "propose",
  "allowed_scopes": [
    "memory_correction",
    "prompt_patch",
    "context_playbook",
    "evaluation_fixture"
  ],
  "forbidden_scopes": [
    "production_activation",
    "permission_change",
    "credential_change",
    "core_source",
    "base_model_weights",
    "safety_policy",
    "termination_policy",
  ]
}
```

```
    "telemetry_retention",
    "gate_thresholds",
    "holdout_sets",
    "corrigibility_invariants",
    "isolation_tier",
    "regression_fixture_removal"
  ],
  "auto_promote": false,
  "requires_independent_evaluator": true,
  "requires_human_approval": true,
  "requires_verifier": true,
  "requires_full_regression_pass": true,
  "requires_safety_suite_pass": true,
  "requires_statistical_plan": "evals/analysis_plan.json",
  "reward_hacking_detectors": [
    "length_exploit",
    "judge_mimicry",
    "fixture_memorization",
    "refusal_gaming",
    "format_exploit",
    "golden_task_divergence"
  ],
  "canary": {
    "strategy": "group_sequential",
    "max_traffic_pct": 5,
    "guardrails": [
      "success_rate",
      "cpst",
      "p95_job_ms",
      "safety_block_rate",
      "unsupported_claim_rate"
    ]
  },
  "rollback_required": true,
  "ledger_ref": "improvement/ledger.json"
}
```

19. Operator and host APIs

All v2 routes are preserved under `/api/v3/...`

Method	Path	Purpose
GET	<code>/api/v3/agents/{id}/structure</code>	Structure and schema
GET	<code>/api/v3/agents/{id}/resolved</code>	MRO, capabilities, locks
POST	<code>/api/v3/agents/{id}/compose-preview</code>	Full dry run

Method	Path	Purpose
GET	/api/v3/agents/{id}/runtime/plan	Compiled DAG
GET	/api/v3/agents/{id}/runtime/capabilities	Runtime capability view
GET	/api/v3/agents/{id}/capabilities/matrix	Asserted, verified, refuted
POST	/api/v3/agents/{id}/capabilities/verify	Run conformance
GET	/api/v3/agents/{id}/runtime/context-budget	Segment budget and use
GET	/api/v3/agents/{id}/cache/stats	Per-tier statistics
POST	/api/v3/agents/{id}/cache/invalidate	Audited scoped invalidation
GET	/api/v3/agents/{id}/protocols	Protocol state
GET	/api/v3/agents/{id}/plugins	Resolved plugins
POST	/api/v3/agents/{id}/plugins/validate	Validate without load
GET	/api/v3/agents/{id}/memory/policy	Memory policy
GET	/api/v3/agents/{id}/memory/hierarchy	Residency and paging
POST	/api/v3/agents/{id}/memory/query	Audited query
POST	/api/v3/agents/{id}/memory/write-candidate	Propose write
POST	/api/v3/agents/{id}/memory/consolidate	Trigger offline job
DELETE	/api/v3/agents/{id}/memory/{memory_id}	Controlled deletion
POST	/api/v3/agents/{id}/memory/{memory_id}/verify-deletion	Run deletion probes
GET	/api/v3/traces/{trace_id}	Operational trace
POST	/api/v3/traces/{trace_id}/replay	Dry replay
POST	/api/v3/traces/{trace_id}/replay?counterfactual=	Counterfactual replay
GET	/api/v3/traces/{trace_id}/root-cause	Failure classification
GET	/api/v3/artifacts/{id}/evidence-graph	Claim attribution
GET	/api/v3/agents/{id}/safety/incidents	Incident records
POST	/api/v3/agents/{id}/safety/redteam	Robustness suite
GET	/api/v3/agents/{id}/improvement/candidates	Candidate list

Method	Path	Purpose
POST	/api/v3/agents/{id}/improvement/candidates/{cid}/evaluate	Sandbox evaluation
POST	/api/v3/agents/{id}/improvement/candidates/{cid}/approve	Human approval
POST	/api/v3/agents/{id}/improvement/rollback/{version}	Rollback
GET	/api/v3/agents/{id}/improvement/ledger	Improvement ledger
GET	/api/v3/agents/{id}/regression/suite	Fixture inventory
GET	/api/v3/agents/{id}/corrigibility/attestation	Invariant attestation
GET	/api/v3/agents/{id}/validation/report	Full report

All mutations require:

- authenticated actor;
- reason;
- expected parent version;
- dry-run response;
- append-only audit;
- explicit human approval for production effects.

No agent-identity endpoint can write corrigibility invariants.

20. Consolidated error catalogue

20.1 Inheritance, skills, and identity

Code	Condition	Default action
INH_CYCLE	Parent graph cycle	Abort
INH_PARENT_MISSING	Parent folder or required file missing	Abort
INH_STRUCTURE_MISMATCH	Incompatible family	Abort
INH_DEPTH	Depth exceeds limit	Abort
INH_PARENT_LIMIT	Parent limit exceeded	Abort
INH_SURFACE_UNKNOWN	Unknown inherited surface	Abort
INH_SELF_PARENT	Child names itself	Abort
INH_RESOLVED_DRIFT	Generated resolution differs from source hashes	Abort
INH_FIXTURE_REMOVAL	Fixture removed without waiver	Abort
SKL_TOGGLE_UNKNOWN	Toggle targets unknown skill	Abort
SKL_MISSING_FILES	Enabled skill lacks required files	Abort

Code	Condition	Default action
SKL_TOOL_LEAK	Skill attempts undeclared tool grant	Abort
IDN_NAMED_PERSON	Named-person overlay lacks approval	Abort
IDN_LICENSE_CLAIM	Persona asserts a real license	Abort
IDN_DISCLOSURE_MISSING	Overlay lacks disclosure	Abort
GATE_NETWORK	Inheritance or identity attempts network grant	Abort
GATE_ACTIVATION	Agent attempts production activation	Abort

20.2 Performance, cache, and context

Code	Condition	Default action
PERF_PLAN_CYCLE	DAG cycle	Abort
PERF_UNSAFE_PARALLELISM	Unsafe side-effect concurrency	Abort
PERF_BUDGET_EXCEEDED	Compile-time budget infeasible	Abort or escalate
PERF_DEADLINE	Deadline exceeded	Cancel with bounded failure
PERF_CACHE_SCOPE	Cache crosses boundary	Abort and purge
PERF_ROUTE_UNAVAILABLE	Route unavailable	Approved fallback only
PERF_STOP_RULE	Stopping rule absent or unauditable	Abort adaptive mode
PERF_SPECULATION_COMMIT	Speculation commits before guard	Abort and compensate
PERF_FALLBACK_INVALID	Fallback lacks validation	Containment stop
CACHE_INVALIDATION	Dependency not invalidated before read	Disable tier and purge
CACHE_EQUIVALENCE	Cached and baseline semantics diverge	Disable tier
CACHE_SEMANTIC_REUSE	T3 exceeds false-reuse threshold	Disable T3
CTX_PRESERVATION	Compaction loses invariant or critical constraint	Escalate or stop
CTX_BUDGET	Required pinned context exceeds budget	Reject or split task

20.3 Compatibility and observability

Code	Condition	Default action
CMP_PROTOCOL_VERSION	Unsupported protocol revision	Abort
CMP_CAPABILITY_MISSING	Required capability absent	Abort

Code	Condition	Default action
CMP_ASSERTED_UNVERIFIEDED	Production binding targets unverified capability	Abort
CMP_CAPABILITY_DRIFT	Verified capability later fails	Quarantine route
CMP_SCHEMA_INCOMPATIBLE	Schema major mismatch	Abort
CMP_TRACE_CONTEXT	Trace/deadline/auth/taint lost	Abort high-risk exchange
CMP_SEMCONV_VERSION	Semconv changes	Quarantine export mapping
CMP_TEMPLATE_DRIFT	Template digest changes	Re-run conformance
CMP_TOKENIZER_DRIFT	Tokenizer digest changes	Re-run conformance
CMP_JSON_SCHEMA_PROFILE	Unsupported schema construct	Fail compose
OBS_TRACE_BROKEN	Root or parent span missing	Abort high-risk run
OBS_REDACTION	Content cannot be safely redacted	Metadata only or abort
OBS_COT_EXPORT	Hidden reasoning targeted for export	Block and alert
OBS_AUDIT_UNAVAILABLE	Exporter and local audit unavailable	Containment stop
OBS_SAMPLING_LOSS	Mandatory trace category would be dropped	Containment stop
OBS_EVIDENCE_GRAPH	Required graph missing or invalid	Block artifact
OBS_UNSUPPORTED_CLAIM	Unsupported claim enters prohibited path	Block write or artifact

20.4 Plugins and memory

Code	Condition	Default action
PLG_MANIFEST_INVALID	Invalid manifest	Abort load
PLG_SIGNATURE	Signature missing or invalid	Abort load
PLG_PERMISSION	Undeclared access	Kill plugin and abort node
PLG_ABI	ABI incompatibility	Abort load
PLG_DEPENDENCY_CYCLE	Dependency cycle	Abort load
PLG_ISOLATION_TIER	Tier absent or too weak	Abort load
PLG_SBOM_MISSING	Production SBOM absent	Abort load
PLG_SUPPLY_CHAIN	Provenance unverifiable	Abort load
PLG_SCAN	Vulnerability threshold exceeded	Abort load
PLG_HANDLE_FORGERY	Forged or unauthorized handle	Kill plugin and alert
PLG_HOT_SWAP	Replacement fails shadow validation	Keep prior version

Code	Condition	Default action
MEM_PROVENANCE	Write lacks valid provenance	Quarantine
MEM_SCOPE	Tenant or subject mismatch	Abort
MEM_CONFLICT	Irreconcilable memory conflict	Abstain or escalate
MEM_POISON	Suspected malicious memory	Quarantine
MEM_TRUST_TIER	Trust promotion violates policy	Reject
MEM_TAINT_LAUNDER	Transform loses taint	Reject output
MEM_PAGE_LOSS	Page-out loses retrievable state	Abort
MEM_DELETE_INCOMPLETE	Derived copy remains	Fail deletion SLA
MEM_UNLEARN_VERIFY	Post-deletion probe retrieves content	Fail deletion SLA

20.5 Improvement, safety, corrigibility, and validation

Code	Condition	Default action
IMP_SCOPE	Candidate touches forbidden scope	Reject and alert
IMP_HOLDOUT_LEAK	Held-out leakage	Reject
IMP_REGRESSION	Release metric regression	Reject or rollback
IMP_UNSIGNED	Missing approved signature	Reject
IMP_SELF_APPROVAL	Self-promotion attempt	Reject and alert
IMP_ROLLBACK	Rollback absent or untested	Block deployment
IMP_REWARD_HACK	Reward-hacking detector fails	Reject
IMP_STAT_UNDERPOWERED	Power requirement not met	Block promotion
IMP_VERIFIER_MISSING	Objective lacks verifier	Reject objective
IMP_VERIFIER_DEPENDENCE	Verifier independence fails	Reject
IMP_CORRIGIBILITY	Invariant mismatch or attempted mutation	Containment stop
IMP_LEDGER_WRITE	Ledger integrity failure	Block promotion
SAF_INJECTION	Injection detected	Block affected instruction/action
SAF_HIJACK	Goal or tool hijack detected	Block and incident
SAF_EXFILTRATION	Secret or PII egress attempt	Block and incident
SAF_TAINT	Taint integrity failure	Block
SAF_TERMINATION	Guard fails to halt correctly	Containment stop

Code	Condition	Default action
SAF_EXTERNAL_EFFECT	Unapproved external action	Block and incident
SAF_CASCADE	Cross-agent escalation	Halt exchange graph
VAL_PLAN_DRIFT	Analysis plan changes after run start	Invalidate run
VAL_FREEZE_DRIFT	Frozen comparison input changes	Invalidate comparison
VAL_EXCLUSION	Undeclared result exclusion	Invalidate report
CIT_UNVERIFIED	Reference lacks accepted audit	Block release
CIT_MISMATCH	Identifier/title/venue mismatch	Delete or correct reference
CIT_NUMERIC_CLAIM	Numeric claim not located	Withdraw claim

21. Validation specification, harness, and report

21.1 Honesty classes

Class	Meaning
MEASURED_LOCAL	Executed on the CASOPS implementation
MEASURED_EXTERNAL	Reported by an audited external source
STATIC_PASS	Verified from specification or schemas
NOT_RUN	Requires an unsupplied implementation
BLOCKED	Release cannot proceed

21.2 Measurement statement

Delivered:

- measurable targets;
- normative definitions;
- powered statistical procedure;
- fixture layout;
- CLI contract;
- report schema;
- failure semantics.

Not delivered:

- local runtime numbers;
- production certification;
- a cleared citation audit.

All figures labelled TARGET are thresholds, not observed outcomes.

21.3 Harness

```

evals/
  analysis_plan.json
  benchmarks.json
  baselines.json
  fixtures/
    perf/{parallel_tool,cache_equivalence,context_rot,kill_switch}/

compat/{model_profiles,mcp,a2a,cloudevents,trace_context,semconv,template_drift}/
obs/{fault_injection,redaction,replay,sampling,evidence_graph}/
plugins/{zero_core_change,permission_denial,isolation,supply_chain,abi}/
memory/{longmem_profile,update,forget,poison,deletion,tce}/
improve/{holdout,reward_hacking,canary_sim,rollback}/
safety/{indirect_injection,hijack,exfiltration,termination,taint_laundering}/
corrigibility/{inv01..inv12_negative}/
regression/
reports/
  <iso8601>-<compose_hash>/
  report.json
  raw/
  statistics.json
  citation-audit.json

```

Invocation:

```

casops-eval run \
  --agent agents/video.showrunner \
  --baseline evals/baselines.json#v2-frozen \
  --plan evals/analysis_plan.json \
  --suite perf,compat,obs,plugins,memory,improve,safety,corrigibility \
  --arms baseline,candidate \
  --paired \
  --seed 20260824 \
  --out evals/reports/

```

The tool:

- exits non-zero on any blocking gate;
- records the pre-run plan digest;
- refuses a valid report if the plan changes after start;
- retains all raw rows;
- counts timeouts and errors as failures unless the estimand explicitly states otherwise.

21.4 Statistical protocol

21.4.1 Freeze list

Freeze:

- dataset and hash;
- child and parent hashes;
- model and adapter revisions;
- tokenizer and template digests;
- quantization;
- hardware;
- framework;
- tool fixtures;
- memory seed;
- cache mode;
- seed where verified;
- retry and timeout policy;
- network conditions;
- evaluator;
- semconv schema;
- analysis-plan digest.

21.4.2 Design

- identical task set across arms;
- task as a blocking factor;
- randomized interleaving;
- cold and warm cache reported separately;
- no undocumented exclusions;
- declared primary estimand for every gate;
- stratification declared before execution;
- effect sizes and interval estimates reported.

21.4.3 Prospective power and sample size

A fixed sample count does not guarantee power.

For every inferential gate, `analysis_plan.json` declares:

- baseline or nuisance-rate assumptions;
- minimum effect or non-inferiority/equivalence margin;
- alpha;
- target power;
- paired discordance assumptions where binary;
- expected attrition or invalid-run rate;
- calculation or simulation method;
- maximum and minimum sample size;
- allowed blinded sample-size re-estimation.

Default target power:

- at least 80% for ordinary performance and quality gates;

- at least 90% for high-risk safety, deletion, capability, or promotion gates where inference—not a full deterministic fixture—is used.

Minimum floors:

Claim	Minimum floor	Final n
p50 latency	300 per arm	Maximum of floor and powered n
p95 latency	300 per arm	Maximum of floor and powered n
p99 latency	1000 per arm	Otherwise labelled indicative
Binary task success	400 paired tasks	Maximum of 400 and prospectively powered n
CPST	300 paired tasks	Maximum of floor and powered n
Memory rates	400 paired tasks	Maximum of 400 and powered n
Safety attack rate	Full declared suite	No sampling; exact interval reported

For paired binary outcomes, power calculations use expected discordant-pair probabilities, not an unpaired approximation alone. Historical or pilot estimates may initialize nuisance parameters. Any re-estimation must be blinded to treatment labels or pre-specified as group-sequential.

An underpowered result is not a pass, even if its point estimate exceeds the threshold.

21.4.4 Tests

Superiority

- one-sided paired test at the pre-registered alpha;
- paired McNemar or a declared paired risk-difference procedure for binary outcomes;
- paired bootstrap, permutation, or model-based procedure for continuous/skewed metrics;
- effect size and confidence interval required.

Non-inferiority

- one-sided NI procedure against a declared margin;
- pass only when the appropriate confidence bound excludes inferiority beyond the margin;
- default one-sided alpha is 0.025 for release-critical quality preservation unless the analysis plan justifies another value;
- “not statistically different” is not evidence of non-inferiority.

Equivalence

- TOST is used only when both lower and upper equivalence bounds are material;
- at $\alpha=0.05$ per one-sided test, report the corresponding 90% confidence interval;
- cache-on/off equivalence may use TOST when the estimand and margin are suitable.

Multiple claims

- every gate is reported;
- no favorable subset may be selected after observation;

- if a family-wide superiority claim is made, a pre-specified multiplicity method such as Holm control is required;
- independent operational gates may remain individually blocking without being merged into one omnibus claim.

Canary monitoring

- uses pre-registered group-sequential boundaries;
- declares looks, spacing, alpha spending, and stop rules;
- naive repeated peeking is prohibited.

21.4.5 Determinism

Token-identical replay is gated only when `batch_invariant_kernels` is verified.

Otherwise replay equivalence means:

- identical external observations;
- identical tool and memory fixtures;
- equivalent route and policy decisions;
- identical validator verdicts within declared tolerance.

Claiming token-level determinism without verified batch invariance is prohibited.

21.5 Release gates

21.5.1 Performance targets against frozen v2

Satisfy A or B.

Gate	Requirement
A — efficiency	p95 improves $\geq 35\%$; CPST improves $\geq 30\%$; task success is non-inferior within 1pp
B — quality	Task success improves $\geq 5\text{pp}$ with superiority CI excluding zero; p95 and CPST regress $\leq 10\%$ unless separately approved

Additional targets:

Check	Threshold
Parallel three-tool fixture	$\text{max}(\text{tool durations}) + 20\%$
Warm-cache TTFO	$\geq 25\%$ improvement over cold
Cache equivalence	Within declared two-sided equivalence margin
T3 false reuse	$\leq 0.5\%$
Cache scope	Zero violations
CPE	≥ 0.70

Check	Threshold
CRR	≥ 0.40
Goodput	$\geq 25\%$ improvement at fixed accelerator count
Stopping rule	$\geq 30\%$ refinement-cost reduction with success non-inferior within 1pp
Optimizer kill switches	100% return to validated baseline
Mandatory control failure	100% containment-stop; zero bypass
Context rot	Non-inferior to oracle-short context within 3pp

21.5.2 Compatibility

- 100% mandatory model contract tests;
- at least three adapter profiles plus deterministic test adapter;
- all production capabilities verified;
- no asserted-unverified binding;
- 100% injected capability drift detected;
- 100% tokenizer/template drift detected;
- pinned and prior MCP revisions negotiate;
- unknown major versions fail closed;
- extensions remain inert unless allow-listed;
- A2A preserves taint, deadline, trace, and authorization;
- 100% CloudEvents schema validation;
- 100% W3C trace continuity;
- 100% alias coverage for gate-bearing telemetry fields.

21.5.3 Observability

- exactly one root trace per run;
- $\geq 99.9\%$ valid span relationships;
- 100% representation of tool, plugin, memory-write, peer, policy, safety, promotion, and rollback actions;
- no raw CoT export;
- no reasoning-monitor leak;
- 100% secret-redaction fixture pass;
- $RCA@1 \geq 85\%$;
- replay equivalence $\geq 95\%$ at the applicable level;
- evidence graph for every claim-bearing artifact;
- unsupported-claim rate $\leq 1\%$;
- mandatory tail retention survives induced budget exhaustion.

21.5.4 Extensibility

- tool, modality, and evaluator plugin installed with zero core changes;
- all undeclared permission attempts denied;
- all handle-forgery and unauthorized delegation attempts denied;
- $I1 \leq 1$ ms median and $\leq 3\%$ p95;

- $I2 \leq 5\%$ p95;
- $I3 \leq 15\%$ p95;
- plugin removal leaves no unresolved capability;
- invalid digest, signature, ABI, schema, SBOM, provenance, or scan fails closed;
- unauthorized tier downgrade blocked;
- regressing hot swap blocked.

21.5.5 Memory

Check	Threshold
Profile score	$\geq 12\text{pp}$ macro or $\geq 25\%$ relative improvement
Evaluation validity	Contamination check and domain golden-task confirmation
Memory prompt tokens	$\geq 50\%$ reduction
TCE	$\geq 30\%$ improvement
Unsupported memory answer rate	$\leq 1\%$
Update and selective forgetting	$\geq 97\%$
Staleness rate	$\leq 2\%$
Provenance	100%
Trust integrity	No T3 used as factual support
DCR	100% by post-deletion probe
Cross-tenant retrieval	Zero
MPR	$\geq 95\%$; no poison reaches T0/T1

21.5.6 Improvement

A promoted candidate must:

- improve held-out success by at least 5pp with superiority evidence, or reduce CPST by at least 10% with quality non-inferior within the declared margin;
- meet prospectively calculated power;
- pass the full regression and safety suites;
- introduce no compatibility regression;
- preserve mandatory telemetry;
- use isolated held-out data;
- pass all reward-hacking detectors;
- carry an independent verifier attestation;
- pass group-sequential canary guardrails;
- receive independent human approval;
- be signed;
- pass rollback RTO;
- append a complete ledger entry;

- preserve corrigibility.
- There must be zero successful self-promotions. Promotion-induced regression rate must be at most 2% over a rolling window of at least 20 promotions.

21.5.7 Safety and corrigibility

- all \$14.4 gates pass;
- INV-01 through INV-12 negative fixtures abort correctly;
- 100% compose attestation coverage;
- cancellation succeeds within deadline at all isolation tiers;
- mandatory-control unavailability always containment-stops;
- no bypass mode exists.

21.6 Citation audit

Release requires:

- zero [D], [C], or [K] markers;
- matching titles, identifiers, venues, and dates;
- every numeric claim located;
- no future-dated verification;
- committed citation-audit.json.

21.7 Static report

Domain	Finding	Status
Date integrity	Future date removed	STATIC_PASS
v2 defects	DEF-001 through DEF-004 addressed	STATIC_PASS
v3 draft defects	DEF-005 through DEF-008 addressed	STATIC_PASS
Architecture	Nine planes and boundaries defined	STATIC_PASS
Performance	DAG, admission, goodput, compute, switches, gates	STATIC_PASS
Cache/context	Keying, invalidation, equivalence, lifecycle	STATIC_PASS
Compatibility	Verified capabilities, drift, pinning, profiles	STATIC_PASS
Observability	Evidence graph, sampling, RCA, bounded monitor	STATIC_PASS
Plugins	Isolation, object capabilities, supply chain, ABI	STATIC_PASS
Memory	Trust, hierarchy, poisoning, consolidation, deletion	STATIC_PASS
Improvement	Verifiers, reward checks, regression ratchet, ledger	STATIC_PASS
Safety	Taint, termination, zero-tolerance categories	STATIC_PASS
Corrigibility	Twelve invariants and negative fixtures	STATIC_PASS
Statistics	Power-derived n; NI/equivalence separated	STATIC_PASS

Domain	Finding	Status
Standalone completeness	Retained requirements reproduced	STATIC_PASS
Error catalogue	Consolidated \$20 supplied	STATIC_PASS
Harness	Layout, CLI, reports, exit semantics	STATIC_PASS
Local performance	Runtime not supplied	NOT_RUN
Local compatibility	Adapters not supplied	NOT_RUN
Local observability	Collector not supplied	NOT_RUN
Local plugins	Runtime not supplied	NOT_RUN
Local memory	Backend not supplied	NOT_RUN
Local improvement	Trainer not supplied	NOT_RUN
Local safety/corrigibility	Runtime not supplied	NOT_RUN
Citation audit	No accepted audit artifact supplied	BLOCKED
Production certification	Requires all gates	BLOCKED

21.8 External evidence retained

These are external claims, not CASOPS results, and remain citation-gated.

Pattern	Reported external result	Class
Paged KV attention	2–4× serving throughput at comparable latency	MEASURED_EXTERNAL, E1 [C]
Structured generation / RadixAttention	Up to 6.4× throughput	MEASURED_EXTERNAL, E1 [C]
Parallel function DAG	Up to 3.7× latency, 6.7× cost, ~9% accuracy	MEASURED_EXTERNAL, E2 [C]
Learned model routing	More than 2× cost reduction in some settings	MEASURED_EXTERNAL, E2 [C]
Workflow search	5.7% average over evaluated baselines	MEASURED_EXTERNAL, E2 [C]
Agent-aware KV management	Reported agent-serving gains	MEASURED_EXTERNAL, E3 [D]
Policy-driven agentic serving	Reported serving gains	MEASURED_EXTERNAL, E3 [D]
Multi-agent workload-aware cache	Reported reuse gains	MEASURED_EXTERNAL, E3 [D]

Pattern	Reported external result	Class
Graph/PPR memory	Up to 20% multi-hop improvement	MEASURED_EXTERNAL, E1 [C]
Graph memory v2	7% associative-memory improvement	MEASURED_EXTERNAL, E2 [C]
Extraction memory	Reported latency and token savings	MEASURED_EXTERNAL, E3 [C]
Self-refinement	~20% average absolute improvement	MEASURED_EXTERNAL, E1 [C]
Reflective context evolution	Reported lower-rollout gains	MEASURED_EXTERNAL, E2/E3 [K]
Self-editing coding agents	Reported coding gains	MEASURED_EXTERNAL, E4 [C]
Agent Lightning + 14.6	Withdrawn; identifier/claim unresolved	—

21.9 Conclusion

Item	Verdict
Specification completeness	PASS
Architecture quality	PASS
Date integrity	PASS
Quantitative criteria	PASS
Statistical protocol	PASS
Harness contract	PASS
Research traceability	BLOCKED pending audit
Executed implementation validation	NOT RUN
Specification status	DRAFT
Deployment recommendation	NO-GO

22. Migration from v2

22.1 Defaults

cache.tiers	= [T0]
context.compaction	= disabled
compute_controller.mode	= fixed

```

capability_verification      = required
memory.hierarchy             = flat H1
memory.consolidation         = disabled
memory.security.trust_tiers  = enabled
plugins.isolation            = I2 minimum for existing plugins
safety.plane                 = enabled
safety.termination           = enforced
corrigibility.invariants     = enforced
observability.sampling        = tail with mandatory retention
observability.evidence_graph = enabled for claim-bearing artifacts
improvement.mode             = inherited but capped at propose

```

Safety, termination, corrigibility, and capability verification activate unconditionally. A refuted existing capability is a discovered latent defect, not a migration regression.

22.2 Steps

1. Copy the v2 folder.
2. Set schema to 3.0 and structure to `casops.common_agent.v3`.
3. Install host-owned corrigibility invariants.
4. Add mandatory safety and termination policy.
5. Add required v3 directories and defaults.
6. Verify every capability.
7. Resolve refutations.
8. Pin tokenizer, template, semconv, protocol, and adapter revisions.
9. Assign isolation and collect SBOM/provenance.
10. Seed regression fixtures from known failures.
11. Author the analysis plan before measurement.
12. Generate locks.
13. Run v2/v3 golden-envelope comparison.
14. Verify no unauthorized tool, network, identity, permission, or production change.
15. Freeze the powered v2 baseline.
16. Enable one optional v3 feature at a time.
17. Run its gates.
18. Complete citation audit.
19. Record migration report.
20. Promote only after all mandatory gates pass.

22.3 Backward compatibility

- v1/v2 prompts, rubrics, skills, identity, and parents remain readable.
- v3-only parents cannot enter a v2 child without explicit down-conversion.
- Older consumers may ignore namespaced v3 metadata.
- Safety, provenance, and corrigibility fields cannot silently disappear during down-conversion.
- Regression and safety fixtures never down-convert away.

23. Traceability

Need	Requirements	Acceptance
Latency and utilization	FR-PERF-001–017, 101–110	§21.5.1
Cache correctness	FR-CACHE-001–009	Equivalence and scope fixtures
Context lifecycle	FR-CTX-001–007	Context-rot gate
Adaptive compute	FR-PERF-103/104	Stopping-rule gate
Model interoperability	FR-CMP-001–010, 101–107	Verified matrix
Telemetry stability	FR-CMP-108–111	Alias and version fixtures
Tool interoperability	FR-CMP-112–117	MCP conformance
Agent interoperability	FR-CMP-118–121	Peer and taint conformance
Decision provenance	§10	Evidence and RCA gates
Reasoning monitor	FR-OBS-101–105	Zero-leak fixtures
Trace cost	FR-OBS-111–115	Budget-exhaustion fixture
Plugin extensibility	FR-PLG-001–010, 101–118	Plugin gates
Supply chain	FR-PLG-109–113	Fail-closed fixtures
Memory	§12	Memory profile
Deletion	FR-MEM-116–120	DCR and probes
Memory attack resistance	FR-MEM-106, §14.4	MPR
Improvement	§13	Candidate gates
Reward integrity	FR-IMP-101–106	Detector suite
Capability retention	FR-IMP-107–110	Full regression
Safety	FR-SAF-001–012	§21.5.7
Corrigibility	INV-01–12, FR-COR-001–006	Negative fixtures
Skills	FR-SKL-001–010	Compose preview
Persona isolation	FR-IDN-001–012	Disclosure and policy tests
Reproducibility	Locks and replay	Applicable replay gate
Statistical validity	§21.4	Plan and power check
Citation integrity	P29, CIT-GATE-001/002	Audit report
Error consistency	§20	Error-code schema validation

24. Open risks

Risk	Required mitigation
Optimizer improves latency but harms success	CPST plus NI quality gate
Router drift	Shadowing, bounded updates, quarantine
Cache serves stale policy	Invalidate-before-read and approval epochs
Cache crosses boundaries	Full-scope keys, purge, zero-tolerance gate
Compaction loses constraints	Pinned invariants and preservation verifier
Vendor overstates capabilities	Verified-not-asserted binding
Template/tokenizer changes semantics	Digest pinning and re-conformance
External telemetry schema changes	Stable <code>casops.*</code> aliases
Protocol revision breaks semantics	Pinning and dual-revision support
Plugin supply-chain compromise	SBOM, provenance, scan, signature, isolation
Isolation under-assigned	Threat-model minimum and expiring waiver
Tool output injects instructions	Taint and no instruction authority
Taint laundering	Propagation through every transform
Telemetry leaks content	Metadata default and redaction
Trace budget drops critical events	Mandatory tail retention
Monitor channel becomes export back door	Verdict-only output and zero-leak tests
Memory accumulates false facts	Trust tiers and provenance
Agent output hardens into fact	Evidence graph and T3 advisory treatment
Memory becomes stale	Bitemporal records and staleness gate
Deletion leaves derivatives	Tombstones and post-deletion probes
Weight-level influence survives deletion	Explicit limitation and retraining review
Public memory benchmark is misleading	Contamination and domain-golden confirmation
Improvement overfits	Isolated held-out and rotating fixtures
Self-judge rewards its style	Independent verifier
Metric improves while capability regresses	Full regression suite

Risk	Required mitigation
Canary peeking inflates false positives	Group-sequential boundaries
Fixed n is underpowered	Prospective paired power calculation
Power assumptions are wrong	Sensitivity analysis and blinded re-estimation
Equivalence confused with NI	Separate procedures in §21.4
Citation supports a false requirement	CIT-GATE-001
Future-dated evidence is represented as complete	CIT-GATE-002
Determinism is overclaimed	Batch-invariance capability gate
Multi-agent failure cascades	Hop, cycle, taint, and shared-budget controls
Parent mixins create a god agent	Parent/depth limits and deny-list union
Added complexity becomes the failure mode	Optional optimizers have baseline kill switches; mandatory controls fail-stop rather than bypass
Corrigibility path appears inside agent folder	Runtime path must be host-owned, immutable, and outside writable capabilities

25. Research references and citation audit

Markers:

- [A] accepted by v3a audit;
- [D] prior draft claimed verification, not accepted by v3a;
- [C] carried from v2;
- [K] unaudited knowledge-derived reference.

All non-[A] entries are blocked by §21.6.

25.1 Prior-draft verification claims [D]

Serving and caching

- *Learning Agent Execution for KV-Cache Management in Agentic Serving*, arXiv:2608.14624 [D]
- *A Policy-Driven Runtime Layer for Agentic LLM Serving*, arXiv:2605.27744 [D]
- *Workload-Aware Caching for Multi-Agent Systems*, arXiv:2607.20495 [D]

Memory

- *A Survey on Memory Mechanisms in the Era of LLMs*, arXiv:2504.15965 [D]
- *Memory in the Age of AI Agents*, arXiv:2512.13564 [D]

- *A Survey of Agent Memory in the Second Half*, arXiv:2602.06052 [D]
- *Agent Memory: Mechanisms, Evaluation, and Emerging Frontiers*, arXiv:2603.07670 [D]
- *Agent Memory Evaluation: Taxonomy and Empirical Analysis of Evaluation and System Limitations*, arXiv:2602.19320 [D]
- *A Survey on the Security of Long-Term Memory in LLM Agents*, arXiv:2604.16548 [D]

Self-evolving agents

- *A Survey of Self-Evolving Agents: What, When, How, and Where to Evolve*, arXiv:2507.21046 [D]
- *Bridging Foundation Models and Lifelong Agentic Systems*, arXiv:2508.07407 [D]

Observability and protocols

- *OpenTelemetry, Semantic Conventions for Generative AI and Inside the LLM Call: GenAI Observability with OpenTelemetry* [D]
- Prior-draft analysis concerning stable versus experimental `gen_ai.*` attributes [D]
- CNCF, *How Jaeger is Evolving to Trace AI Agents with OpenTelemetry* [D]
- Model Context Protocol revisions and *Versioning and Compatibility* [D]

25.2 Carried from v2 [C]

Performance and workflow

- PagedAttention, SOSP 2023 [C]
- SGLang / RadixAttention, NeurIPS 2024 [C]
- LLMCompiler [C]
- RouteLLM [C]
- AFlow, ICLR 2025 [C]
- Automated Design of Agentic Systems, ICLR 2025 [C]
- EAGLE-3 [C]
- Agentic test-time-compute system analysis [C]

Memory

- HippoRAG, NeurIPS 2024 [C]
- HippoRAG 2, ICML 2025 [C]
- LongMemEval, ICLR 2025 [C]
- LoCoMo [C]
- A-MEM [C]
- MIRIX [C]
- Mem0 [C]
- MemoryAgentBench [C]
- Mem2ActBench [C]
- MemGAS [C]

Improvement

- Self-Refine, NeurIPS 2023 [C]
- Reflexion, NeurIPS 2023 [C]

- Voyager, NeurIPS 2023 [C]
- Promptbreeder, ICML 2024 [C]
- Agent Lightning, arXiv:2508.03680 [C]
- Darwin Gödel Machine [C]
- A Self-Improving Coding Agent [C]
- Self-Adapting Language Models / SEAL [C]

Standards and evaluation

- W3C Trace Context [C]
- CloudEvents JSON event format [C]
- Linux Foundation A2A [C]
- Turpin et al., *Language Models Don't Always Say What They Think* [C]
- Lanham et al., *Measuring Faithfulness in Chain-of-Thought Reasoning* [C]
- ToolSandbox [C]
- GAIA [C]
- SWE-bench [C]
- BFCL [C]

Withdrawn: “Agent Lightning v1.0, arXiv:2608.17528, SWE-bench Verified 41.8%→56.4%.” It remains unusable unless independently resolved, and the numeric claim must not be restored merely because a related Agent Lightning paper exists.

25.3 Knowledge-derived references [K]

- MemGPT-style hierarchical or OS-style memory paging [K]
- Memory operating-system architectures [K]
- Temporal knowledge-graph memory, including Zep/Graphiti-style patterns [K]
- Background or sleep-time consolidation compute [K]
- Agentic context engineering and evolving context playbooks [K]
- GEPA-style reflective prompt evolution [K]
- AgentDojo-style injection and hijacking benchmarks [K]
- InjecAgent-style injection benchmarks [K]
- AgentPoison-style memory-poisoning work [K]
- MAST-style multi-agent failure taxonomies [K]
- Batch invariance and inference nondeterminism [K]
- Long-context degradation or context-rot measurements [K]
- Object-capability sandboxing for extensions [K]

If a citation fails, the related control may remain only when it has a documented independent engineering justification.

26. Document control

Item	Value
------	-------

Item	Value
Owner	Host architecture, CASOPS
Revision	v3a
Structure family	casops.common_agent.v3
Schema	3.0
Date	2026-08-31
Supersedes on approval	v2
Replaces as draft	Future-dated unreleased v3 input
Specification completeness	Yes
Release-ready specification	No; citation and local gates remain blocked
Production implementation certified	No
Automatic production activation	No
Automatic tool, plugin, network, or permission grant	No
Automatic candidate promotion	No
Core self-modification	Research-only
Raw chain-of-thought export	Prohibited
Reasoning monitor	Internal-only, verdict-only, short-retention
Default cache	T0 only
Default memory	None until configured
Default improvement	Disabled or propose-only
Safety plane	Mandatory and non-bypassable
Corrigibility	Mandatory, host-owned, agent-unwritable
Optional optimizer failure	Validated baseline fallback
Mandatory-control failure	Containment stop
Statistical protocol	Pre-registered, paired, power-derived, interval-estimated
Non-inferiority	One-sided NI procedure
Equivalence	TOST only when two-sided equivalence is intended
Citation audit	Release-blocking
Public control plane	Existing FastAPI plane only

Item	Value
Normative diagrams	Inline Mermaid
Deployment recommendation	NO-GO until §21 gates pass

Final delivery statement

Delivered: a complete, standalone v3a specification retaining the v3 architecture and requirements while correcting its date, citation status, statistical protocol, kill-switch contradiction, standalone dependency, and missing consolidated error catalogue.

Not delivered: fabricated runtime results, a falsely cleared citation audit, or production certification.

Required next actions:

- 1. execute CIT-GATE-001 and CIT-GATE-002;
- 2. implement casops-eval;
- 3. freeze a powered v2 baseline;
- 4. run every local gate in §21.5;
- 5. retain v3a as DRAFT until both citation and local validation blockers clear.

End of specification.